

Question 1: Exploratory Data Analysis and Preprocessing

Objective

In this question, you will conduct a structured exploratory data analysis (EDA) and preprocess the **Wine Quality Dataset** to prepare it for downstream analysis. This includes computing **summary statistics**, detecting and removing **outliers using Mahalanobis distance**, applying **feature normalization**, and performing **dimensionality reduction with PCA**. Throughout, you will be required to analyze your results and justify your methodological choices.

- **Dataset:** [Wine Quality Dataset](#)
 - **Focus Areas:** Algorithm analysis, data interpretation, and preprocessing strategies.
-

Key Considerations

- **Logical Flow:** The problem walks through **EDA → Cleaning → Transformation → PCA**
- **Focus on Analysis & Justification:** You must **explain your choices** rather than just implement code.
- **Algorithmic Thinking:** Requires **metric development**, use of **Mahalanobis distance**, and **PCA interpretation**.

Question 1: Exploratory Data Analysis and Preprocessing (30 Points)

Part 1: Summary Statistics (5 Points)

Instructions:

1. Load the **Wine Quality Dataset** and inspect its structure (e.g., feature types, missing values, summary statistics).
2. Compute the following descriptive statistics for each feature, both overall and grouped by wine quality rating:
 - Minimum, Maximum
 - Mean, Trimmed Mean (5%)
 - Standard Deviation
 - Skewness, Kurtosis
3. Present results in a **clear table**.
4. Provide a written **interpretation** of what these statistics reveal about the dataset.

Deliverables:

- Code implementation for computing summary statistics.
- A table summarizing computed values.

- A written analysis of key insights.

```
import kagglehub
import pandas as pd

# Download latest version
url_red = kagglehub.dataset_download("uciml/red-wine-quality-cortez-et-al-2009")
url_white = kagglehub.dataset_download("piyushagni5/white-wine-quality")
df_red = pd.read_csv(f"{url_red}/winequality-red.csv", sep=',')
df_white = pd.read_csv(f"{url_white}/winequality-white.csv", sep=';')

df_red['wine_type'] = 'red'
df_white['wine_type'] = 'white'
df_wine = pd.concat([df_red, df_white], axis=0, ignore_index=True)

df_wine.head()

{"summary": "{\n  \"name\": \"df_wine\", \n  \"rows\": 6497, \n  \"fields\": [\n    {\n      \"column\": \"fixed acidity\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 1.296433757799792, \n        \"min\": 3.8, \n        \"max\": 15.9, \n        \"num_unique_values\": 106, \n        \"samples\": [\n          7.15, \n          8.1, \n          7.3\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"volatile acidity\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.1646364740846772, \n        \"min\": 0.08, \n        \"max\": 1.58, \n        \"num_unique_values\": 187, \n        \"samples\": [\n          0.405, \n          0.21, \n          0.695\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"citric acid\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.14531786489759185, \n        \"min\": 0.0, \n        \"max\": 1.66, \n        \"num_unique_values\": 89, \n        \"samples\": [\n          0.1, \n          0.6, \n          0.37\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"residual sugar\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 4.757803743147445, \n        \"min\": 0.6, \n        \"max\": 65.8, \n        \"num_unique_values\": 316, \n        \"samples\": [\n          18.95, \n          3.2, \n          9.3\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"chlorides\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.03503360137245906, \n        \"min\": 0.009, \n        \"max\": 0.611, \n        \"num_unique_values\": 224, \n        \"samples\": [\n          0.089, \n          0.096, \n          0.422\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"free sulfur dioxide\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 17.74939977200255, \n        \"min\": 1.0, \n        \"max\": 289.0, \n
```

```

{"num_unique_values": 135,\n      "samples": [\n          77.5,\n          65.0,\n          128.0\n      ],\n      "semantic_type": \"\",\n      "description": \"\",\n      "column":\n      \"total sulfur dioxide\",\n      "properties": {\n          "dtype":\n          \"number\",\n          "std": 56.521854522630264,\n          "min":\n          6.0,\n          "max": 440.0,\n          "num_unique_values": 276,\n          "samples": [\n              14.0,\n              149.0,\n              227.0\n          ],\n          "semantic_type": \"\",\n          "description": \"\",\n          "column": \"density\",\n          "properties": {\n              "dtype": \"number\",\n              "std":\n              0.002998673003719041,\n              "min": 0.98711,\n              "max":\n              1.03898,\n              "num_unique_values": 998,\n              "samples": [\n                  0.9918,\n                  0.99412,\n                  0.99484\n              ],\n              "semantic_type": \"\",\n              "description": \"\",\n              "column": \"pH\",\n              "properties": {\n                  "dtype": \"number\",\n                  "std": 0.1607872021039883,\n                  "min": 2.72,\n                  "max": 4.01,\n                  "num_unique_values":\n                  108,\n                  "samples": [\n                      3.74,\n                      3.17,\n                      3.3\n                  ],\n                  "semantic_type": \"\",\n                  "description": \"\",\n                  "column":\n                  \"sulphates\",\n                  "properties": {\n                      "dtype":\n                      \"number\",\n                      "std": 0.14880587361448958,\n                      "min":\n                      0.22,\n                      "max": 2.0,\n                      "num_unique_values": 111,\n                      "samples": [\n                          1.11,\n                          1.56,\n                          0.46\n                      ],\n                      "semantic_type": \"\",\n                      "description": \"\",\n                      "column": \"alcohol\",\n                      "properties": {\n                          "dtype": \"number\",\n                          "std":\n                          1.192711748870997,\n                          "min": 8.0,\n                          "max": 14.9,\n                          "num_unique_values": 112,\n                          "samples": [\n                              11.2,\n                              8.6,\n                              10.5\n                          ],\n                          "semantic_type": \"\",\n                          "description": \"\",\n                          "column":\n                          \"quality\",\n                          "properties": {\n                              "dtype": \"number\",\n                              "std": 0,\n                              "min": 3,\n                              "max": 9,\n                              "num_unique_values": 7,\n                              "samples": [\n                                  5,\n                                  6,\n                                  3\n                              ],\n                              "semantic_type": \"\",\n                              "description": \"\",\n                              "column":\n                              \"wine_type\",\n                              "properties": {\n                                  "dtype":\n                                  \"category\",\n                                  "num_unique_values": 2,\n                                  "samples":\n                                  [\n                                      \"white\",\n                                      \"red\"\n                                  ],\n                                  "semantic_type": \"\",\n                                  "description": \"\"\n                              }\n                          }\n                      }\n                  }\n              }\n          }\n      ],\n      "type": "dataframe",\n      "variable_name": "df_wine"}

```

```
df_wine.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 6497 entries, 0 to 6496
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	fixed acidity	6497 non-null	float64

```

1  volatile acidity      6497 non-null  float64
2  citric acid           6497 non-null  float64
3  residual sugar        6497 non-null  float64
4  chlorides             6497 non-null  float64
5  free sulfur dioxide    6497 non-null  float64
6  total sulfur dioxide   6497 non-null  float64
7  density               6497 non-null  float64
8  pH                   6497 non-null  float64
9  sulphates             6497 non-null  float64
10 alcohol               6497 non-null  float64
11 quality               6497 non-null  int64
12 wine_type            6497 non-null  object

```

```
dtypes: float64(11), int64(1), object(1)
```

```
memory usage: 660.0+ KB
```

```
# --- Function to calculate all desired statistics for a Series ---
```

```
def calculate_descriptive_stats(series):
    stats_dict = {
        'Minimum': series.min(),
        'Maximum': series.max(),
        'Mean': series.mean(),
        'Trimmed Mean (5%)': stats.trim_mean(series.dropna(),
        proportiontocut=0.025), # dropna for safety
        'Standard Deviation': series.std(),
        'Skewness': series.skew(),
        'Kurtosis': series.kurt()
    }
    return pd.Series(stats_dict)
```

```
import numpy as np
from scipy import stats
```

```
numerical_features = [col for col in df_wine.columns if (col !=
'wine_type' and col != 'quality')]
```

```
# Initialize an empty dictionary to store the results
```

```
results =
df_wine[numerical_features].apply(calculate_descriptive_stats)
```

```
# Convert the dictionary of Series into a DataFrame
```

```
stats_df = pd.DataFrame(results)
stats_df = stats_df.reset_index()
stats_df = stats_df.rename(columns={'index': 'Statistic'}) # Rename
'index' column to 'Statistic'
stats_df.set_index('Statistic', inplace=True)
stats_df.round(2)
```

```
{"summary": "{\n  \"name\": \"stats_df\", \n  \"rows\": 7, \n  \"fields\": [\n    {\n      \"column\": \"Statistic\", \n      \"properties\": {\n        \"dtype\": \"string\", \n
```

```

\"num_unique_values\": 7,\n        \"samples\": [\n
\"Minimum\", \n        \"Maximum\", \n        \"Skewness\" \n
n        ], \n        \"semantic_type\": \"\", \n
\"description\": \"\" \n        } \n        }, \n        { \n        \"column\":
\"fixed acidity\", \n        \"properties\": { \n        \"dtype\":
\"number\", \n        \"std\": 4.9489003107472564, \n        \"min\":
1.3, \n        \"max\": 15.9, \n        \"num_unique_values\": 7, \n
\"samples\": [ \n        3.8, \n        15.9, \n        1.72 \n
], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n
} \n        }, \n        { \n        \"column\": \"volatile acidity\", \n
\"properties\": { \n        \"dtype\": \"number\", \n        \"std\":
1.0301433094939938, \n        \"min\": 0.08, \n        \"max\": 2.83, \n
\"num_unique_values\": 7, \n        \"samples\": [ \n        0.08, \n
1.58, \n        1.5 \n        ], \n        \"semantic_type\": \"\", \n
\"description\": \"\" \n        } \n        }, \n        { \n        \"column\":
\"citric acid\", \n        \"properties\": { \n        \"dtype\":
\"number\", \n        \"std\": 0.9055937278934743, \n        \"min\":
0.0, \n        \"max\": 2.4, \n        \"num_unique_values\": 6, \n
\"samples\": [ \n        0.0, \n        1.66, \n        2.4 \n
], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n
} \n        }, \n        { \n        \"column\": \"residual sugar\", \n
\"properties\": { \n        \"dtype\": \"number\", \n        \"std\":
23.572500826982598, \n        \"min\": 0.6, \n        \"max\": 65.8, \n
\"num_unique_values\": 7, \n        \"samples\": [ \n        0.6, \n
65.8, \n        1.44 \n        ], \n        \"semantic_type\": \"\", \n
\"description\": \"\" \n        } \n        }, \n        { \n        \"column\":
\"chlorides\", \n        \"properties\": { \n        \"dtype\":
\"number\", \n        \"std\": 18.951984693153477, \n        \"min\":
0.01, \n        \"max\": 50.9, \n        \"num_unique_values\": 7, \n
\"samples\": [ \n        0.01, \n        0.61, \n        5.4 \n
], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n
} \n        }, \n        { \n        \"column\": \"free sulfur dioxide\", \n
\"properties\": { \n        \"dtype\": \"number\", \n        \"std\":
104.39442988242327, \n        \"min\": 1.0, \n        \"max\": 289.0, \n
\"num_unique_values\": 7, \n        \"samples\": [ \n        1.0, \n
289.0, \n        1.22 \n        ], \n        \"semantic_type\": \"\", \n
\"description\": \"\" \n        } \n        }, \n        { \n        \"column\":
\"total sulfur dioxide\", \n        \"properties\": { \n        \"dtype\":
\"number\", \n        \"std\": 156.37755747213978, \n        \"min\": -
0.37, \n        \"max\": 440.0, \n        \"num_unique_values\": 7, \n
\"samples\": [ \n        6.0, \n        440.0, \n        -0.0 \n
], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n
} \n        }, \n        { \n        \"column\": \"density\", \n
\"properties\": { \n        \"dtype\": \"number\", \n        \"std\":
2.247186600851745, \n        \"min\": 0.0, \n        \"max\": 6.61, \n
\"num_unique_values\": 5, \n        \"samples\": [ \n        1.04, \n
6.61, \n        0.0 \n        ], \n        \"semantic_type\": \"\", \n
\"description\": \"\" \n        } \n        }, \n        { \n        \"column\":
\"pH\", \n        \"properties\": { \n        \"dtype\": \"number\", \n

```

```

\"std\": 1.6415816995174124,\n        \"min\": 0.16,\n        \"max\": 4.01,\n        \"num_unique_values\": 6,\n        \"samples\": [\n2.72,\n        4.01,\n        0.37\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\",\n        \"column\": \"sulphates\",\n        \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 3.033212975110873,\n        \"min\": 0.15,\n        \"max\": 8.65,\n        \"num_unique_values\": 7,\n        \"samples\": [\n        0.22,\n        2.0,\n        1.8\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\",\n        \"column\": \"alcohol\",\n        \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 6.016205891866845,\n        \"min\": -0.53,\n        \"max\": 14.9,\n        \"num_unique_values\": 7,\n        \"samples\": [\n        8.0,\n        14.9,\n        0.57\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n        }\n    }\n}","type":"dataframe"}

```

```

import math
import seaborn as sns
import matplotlib.pyplot as plt

# Plot overall stats
n_cols = 4
n_rows = math.ceil(len(numerical_features) / n_cols)

fig, axes = plt.subplots(n_rows, n_cols, figsize=(n_cols * 6, n_rows * 5))
fig.subplots_adjust(hspace=0.6, wspace=0.4)

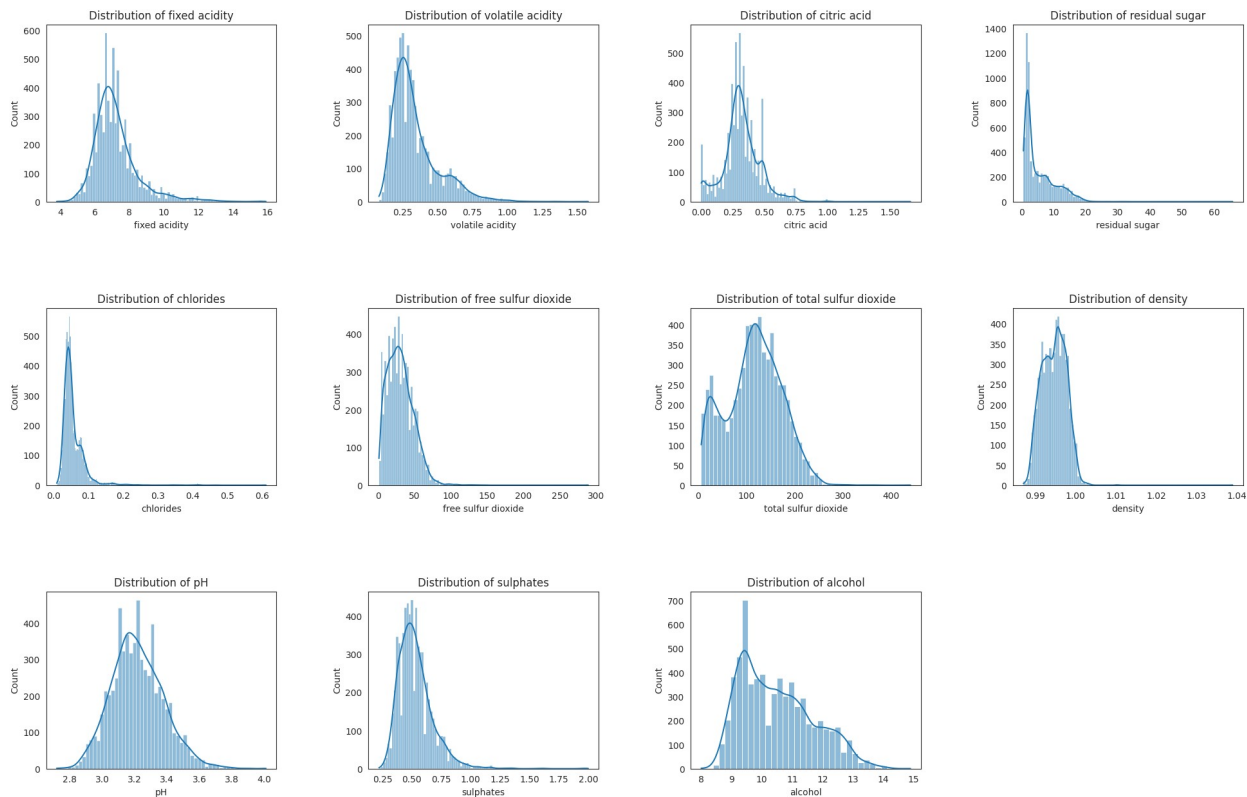
axes = axes.flatten()

for i, col in enumerate(numerical_features):
    ax = axes[i]
    sns.histplot(data=df_wine, x=col, kde=True, ax=ax)

    ax.set_title(f'Distribution of {col}', fontsize=12)
    ax.set_xlabel(col, fontsize=10)
    ax.set_ylabel('Count', fontsize=10)

for j in range(i + 1, len(axes)):
    fig.delaxes(axes[j]) # Delete unused one

```



Structure of dataset

- Total 13 features
- No missing values, each features has total of 6497 rows
- `wine_type` is a categorical feature as it can be either red or white.
- Though `quality` is a numerical feature, it is a categorical as it is integer value in the range of 0 to 10

Overall Dataset Statistics

1. **Highly Skewed Features:** `fixed acidity`, `volatile acidity`, `residual sugar`, `chlorides`, `free sulfur dioxide`, `sulphates`, `pH` show significant positive skewness and often high kurtosis. This indicates the presence of outliers on the higher end for these features. `alcohol` is moderately negatively skewed. `total sulfur dioxide` is notably symmetrical.
2. **Outliers Present:** High kurtosis in `fixed acidity`, `volatile acidity`, `residual sugar`, `chlorides`, `free sulfur dioxide`, `sulphates` suggests that these features have more extreme values (outliers) than expected from a normal distribution. `chlorides` has particularly striking kurtosis, indicating very rare but very high values.
3. **Data Cleaning/Transformation:** For features with high skewness and kurtosis (especially `chlorides`, `residual sugar`, `fixed acidity`), we need to consider following:
 - Investigating the outliers (are they data errors or genuinely rare wines?).

- Applying data transformations (e.g., logarithmic transformation) before building models, as many models assume normally distributed or symmetrical data.
- Impact on Models: The non-normal distributions (especially highly skewed/leptokurtic ones) can impact the performance of statistical models that assume normality.

```
# Stats as per grouped by quality rating
grouped_stats_list = []
for quality_rating, group_df in df_wine.groupby('quality'):
    group_stats =
group_df[numerical_features].apply(calculate_descriptive_stats).round(
2)

    group_stats = group_stats.reset_index()

    group_stats = group_stats.rename(columns={'index': 'Statistic'}) #
Rename 'index' column to 'Statistic'
    group_stats.insert(0, 'Quality', quality_rating) # Add a 'Quality'
column to identify the group

    grouped_stats_list.append(group_stats)

# Concatenate all individual group DataFrames into one
grouped_stats_df = pd.concat(grouped_stats_list)

# Set 'Quality' as the index for a cleaner look
grouped_stats_df.set_index('Quality', inplace=True)

grouped_stats_df

{"summary": "{\n  \"name\": \"grouped_stats_df\", \n  \"rows\": 49, \n
  \"fields\": [\n    {\n      \"column\": \"Quality\", \n
      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\":
2, \n        \"min\": 3, \n        \"max\": 9, \n
        \"num_unique_values\": 7, \n        \"samples\": [\n          3, \n
4, \n          8\n        ], \n        \"semantic_type\": \"\", \n
      \"description\": \"\"\n    }, \n    {\n      \"column\":
\"Statistic\", \n      \"properties\": {\n        \"dtype\":
\"category\", \n        \"num_unique_values\": 7, \n        \"samples\":
[\n          \"Minimum\", \n          \"Maximum\", \n
          \"Skewness\"\n        ], \n        \"semantic_type\": \"\", \n
      \"description\": \"\"\n    }, \n    {\n      \"column\":
\"fixed acidity\", \n      \"properties\": {\n        \"dtype\":
\"number\", \n        \"std\": 4.036318117490026, \n        \"min\":
0.29, \n        \"max\": 15.9, \n        \"num_unique_values\": 45, \n
        \"samples\": [\n          6.6, \n          4.68, \n          15.6\n
        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n
      }, \n    {\n      \"column\": \"volatile acidity\", \n
      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\":
0.8666968041594107, \n        \"min\": -3.06, \n        \"max\": 2.91, \n
        \"num_unique_values\": 38, \n        \"samples\": [\n          0.24, \n
```



```

0.06,\n          1.41\n          ],\n          \"semantic_type\": \"\",\n\"description\": \"\"\n          },\n          {\n          \"column\":  
\"citric acid\", \n          \"properties\": {\n          \"dtype\":  
\"number\", \n          \"std\": 0.9166106043462513, \n          \"min\": -  
1.83, \n          \"max\": 4.74, \n          \"num_unique_values\": 36, \n          \"samples\": [\n          -1.83, \n          0.16, \n          0.03\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\": \"residual sugar\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\":  
10.145228352941757, \n          \"min\": -0.25, \n          \"max\": 65.8, \n          \"num_unique_values\": 44, \n          \"samples\": [\n          -0.25, \n          8.5, \n          0.9\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\":  
\"chlorides\", \n          \"properties\": {\n          \"dtype\":  
\"number\", \n          \"std\": 15.58858443030721, \n          \"min\": -  
2.51, \n          \"max\": 89.21, \n          \"num_unique_values\": 26, \n          \"samples\": [\n          0.05, \n          4.36, \n          0.02\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\": \"free sulfur dioxide\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\":  
51.077594925773084, \n          \"min\": 0.35, \n          \"max\": 289.0, \n          \"num_unique_values\": 44, \n          \"samples\": [\n          3.26, \n          0.65, \n          0.48\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\":  
\"total sulfur dioxide\", \n          \"properties\": {\n          \"dtype\":  
\"number\", \n          \"std\": 101.27997885819788, \n          \"min\": -  
0.92, \n          \"max\": 440.0, \n          \"num_unique_values\": 45, \n          \"samples\": [\n          85.0, \n          -0.6, \n          289.0\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\": \"density\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\":  
2.2811848723347055, \n          \"min\": -1.04, \n          \"max\": 15.59, \n          \"num_unique_values\": 18, \n          \"samples\": [\n          0.99, \n          1.0, \n          -0.26\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\": \"pH\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 1.6358392029656164, \n          \"min\": -1.11, \n          \"max\": 4.01, \n          \"num_unique_values\": 35, \n          \"samples\": [\n          0.0, \n          3.79, \n          3.82\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\":  
\"sulphates\", \n          \"properties\": {\n          \"dtype\":  
\"number\", \n          \"std\": 4.929931118518121, \n          \"min\":  
0.09, \n          \"max\": 32.02, \n          \"num_unique_values\": 40, \n          \"samples\": [\n          0.23, \n          0.14, \n          0.52\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\": \"alcohol\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\":  
5.490374988742282, \n          \"min\": -2.03, \n          \"max\": 14.9, \n

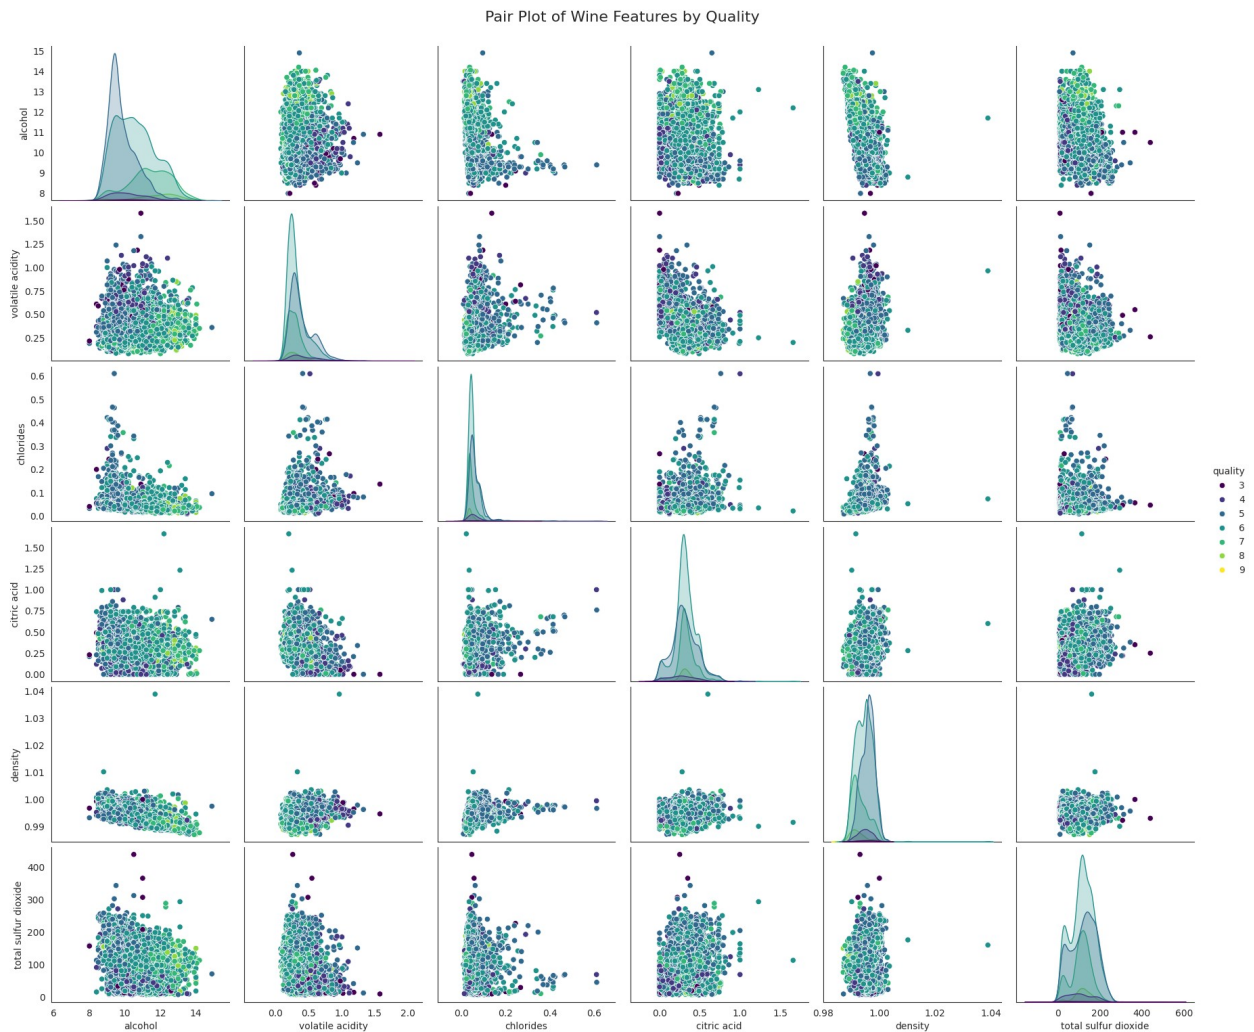
```



```

for continuous variables
sns.pairplot(df_wine[features_for_pairplot], hue='quality',
diag_kind='kde', palette='viridis', height=2.5, aspect=1.2)
plt.suptitle('Pair Plot of Wine Features by Quality', y=1.02,
fontsize=16)
plt.show()

```



Observations

1. The dataset fundamentally contains two distinct groups of wines, which align with their wine_type (red/white) and their quality range.
2. Higher quality wines (mostly white) are generally characterized by higher alcohol, lower volatile acidity, lower chlorides, higher citric acid, and slightly lower density.
3. Lower quality wines (mostly red) exhibit the opposite characteristics: lower alcohol, higher volatile acidity, higher chlorides, lower citric acid, and slightly higher density.

4. Features like `total sulfur dioxide` appear to be less strong discriminators of quality/wine type based on their distribution in this plot.
-

Part 3: Outlier Detection and Removal using Mahalanobis Distance (5 Points)

Instructions:

1. Use the provided **pairwise ellipse plot method** (which builds on the Mahalanobis distance) to assess outliers in the dataset.
2. Select **at least three distinct feature pairs** for visualization.
3. Develop a **numerical outlier metric** based on Mahalanobis distance to systematically identify extreme values.
4. Implement an **algorithm that removes observations identified as outliers** based on this metric.
5. Justify the **choice of threshold** for outlier removal and explain why Mahalanobis distance is appropriate for multivariate data.

Deliverables:

- Code implementing the outlier detection and removal algorithm.
- Pairwise ellipse plots for at least three feature pairs.
- A written explanation of the metric used for outlier detection and removal, including justification of the threshold.

```
import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from numpy.linalg import eigh
import math
from scipy.stats.distributions import chi2
from scipy.stats import norm
from scipy import stats
import seaborn as sns
from sklearn import datasets
from matplotlib import inline

class sigma_ellipse_plot:

    def __init__(self, df=None, target='red',
target_header='wine_type', feature1="", feature2="", std_devs=[1, 2]):

        self.data = df
        self.target = target
        self.feature1 = feature1
        self.feature2 = feature2
```

```

        self.target_header = target_header
        self.std_devs=std_devs
        self.largest_eigenvalue = None
        self.largest_eigenvector = None
        self.smallest_eigenvalue = None
        self.smallest_eigenvector = None
        self.angle = None
        self.mean = None
        self.r_ellipses = None
        self.mu_X = None
        self.mu_Y = None
        self.chisquare_val = None

    def get_data(self):

        self.data = self.data[self.data[self.target_header] ==
self.target].drop(self.target_header, axis =1)[[self.feature1,
self.feature2]]

        return

    def get_eigens(self):

        covariance_matrix = self.data.cov()
        eigenvalues, eigenvectors = eigh(covariance_matrix)

        self.largest_eigenvector =
eigenvectors[np.argmax(eigenvalues)]
        self.largest_eigenvalue = np.max(eigenvalues)
        self.smallest_eigenvector =
eigenvectors[np.argmin(eigenvalues)]
        self.smallest_eigenvalue = np.min(eigenvalues)

        return

    def get_angle(self):

        self.angle = math.atan2(self.largest_eigenvector[1],
self.largest_eigenvector[0])

        return

    def shift_angle(self):

        if self.angle < 0:
            self.angle = self.angle + 2*math.pi

        return

    def get_mean(self):

```

```

        self.mean = self.data.mean()

    return

    def get_chisquare_vals(self):

        self.chisquare_val = []
        for i in range(0, len(self.std_devs)):
            percent_covered = stats.norm.cdf(self.std_devs[i]) -
stats.norm.cdf(self.std_devs[i] * -1)
            self.chisquare_val.append((chi2.ppf(percent_covered,
df=2))*0.5)

        return self.chisquare_val

    def get_ellipses(self):

        chisquare_val = self.get_chisquare_vals()

        self.r_ellipses = []
        for i in range(0, len(self.std_devs)):
            theta_grid = np.linspace(0,2*math.pi, 100)
            phi = self.angle
            self.mu_X = self.mean[0]
            self.mu_Y = self.mean[1]
            a = chisquare_val[i] * math.sqrt(self.largest_eigenvalue)
            b = chisquare_val[i] * math.sqrt(self.smallest_eigenvalue)

            ellipse_x_r = a * np.cos(theta_grid)
            ellipse_y_r = b * np.sin(theta_grid)

            R = [[math.cos(phi), math.sin(phi)], [-math.sin(phi),
math.cos(phi)]]

            ellipses = np.array([ellipse_x_r, ellipse_y_r])

            r_ellipse = ellipses.T.dot(R).T

            self.r_ellipses.append(r_ellipse)

        return

    def get_labels(self, special_phrase=None):

        labels = []
        for i in range(0, len(self.std_devs)):

            if special_phrase is None:
                label = str(self.std_devs[i]) + " std. dev. from mean"
                labels.append(label)
            else:

```

```

        label = special_phrase + str(self.std_devs[i]) + "
std. dev. from mean"
        labels.append(label)

    return labels

def pipeline(self):

    self.get_data()
    self.get_eigens()
    self.get_angle()
    self.shift_angle()
    self.get_mean()
    self.get_ellipses()

    return self.data, self.r_ellipses, self.mu_X, self.mu_Y

# --- Three distinct feature pairs for visualization ---

from matplotlib.patches import Ellipse
feature_pairs = [
    ('alcohol', 'volatile acidity'),
    ('chlorides', 'pH'),
    ('fixed acidity', 'density')
]
# Set confidence level for ellipses (e.g., 95%)
confidence_level_ellipse = 0.95

plt.figure(figsize=(18, 6))

# Create a figure and a set of subplots (1 row, 3 columns)
fig, axes = plt.subplots(1, 3, figsize=(22, 6))
fig.suptitle('Error Ellipses for Red and White Wine Feature Pairs',
fontsize=20, y=1.02)

colors_for_plot_red = ['#1065c0', '#1a8ae5']
colors_for_plot_white = ['#4c9a2a', '#68bb59']

for ax, (feature1, feature2) in zip(axes, feature_pairs):
    custom_std_devs = [1.5, 2.5]

    ## Species Specific Ellipse Generation ##
    red_ellipses_obj = sigma_ellipse_plot(df=df_wine, target='red',
feature1=feature1, feature2=feature2, std_devs=custom_std_devs) # Use
custom_std_devs
    red_df, red_ellipses, red_mu_X, red_mu_Y =
red_ellipses_obj.pipeline()
    red_plot_labels = red_ellipses_obj.get_labels()

    white_ellipses_obj = sigma_ellipse_plot(df=df_wine, target='white',
feature1=feature1, feature2=feature2, std_devs=custom_std_devs) # Use

```



```

custom_std_devs
white_df, white_ellipses, white_mu_X, white_mu_Y =
white_ellipses_obj.pipeline()
white_plot_labels = white_ellipses_obj.get_labels()

## Visualization of Ellipse Plots ##
sns.set_style("white")
df_subset = df_wine[[feature1, feature2, 'wine_type']]

# palette order - white, red
sns.scatterplot(data=df_subset, x=feature1, y=feature2,
hue="wine_type", palette = ["#0747a1", '#d24e01'], ax=ax)

# Set titles and legends for each subplot
ax.set_title(f"{feature1.title()} vs. {feature2.title()} ",
fontsize=14)
ax.legend(loc='best')
ax.grid(True, linestyle='--', alpha=0.6)

ax.scatter(red_mu_X, red_mu_Y, c='#1f51ff', s=150, label = 'Red Wine
mean', marker='$\mu$')
ax.scatter(white_mu_X, white_mu_Y, c='#ff6700', s=150, label =
'White Wine mean', marker='$\mu$')


ax.legend()

# Plots the ellipses for each wine type, with the appropriate colors
for i in range(0, len(red_ellipses)):
    ax.plot(red_ellipses[i][0] + red_mu_X, red_ellipses[i][1] +
red_mu_Y, colors_for_plot_red[i], label= red_plot_labels[i])
    ax.plot(white_ellipses[i][0] + white_mu_X, white_ellipses[i][1]
+ white_mu_Y, colors_for_plot_white[i], label= white_plot_labels[i])

```

<Figure size 1800x600 with 0 Axes>

Error Ellipses for Red and White Wine Feature Pairs



Interpretation of Ellipse Plots:

- We clearly see two distinct clusters of points, one for 'red' wines and one for 'white' wines.
- Each cluster have its own ellipse, tightly fitting around its respective data points.
- Points that fall outside their respective wine type's ellipse are candidates for outliers within that specific wine type.

Numerical Outlier Metric (Mahalanobis Distance) and Removal Algorithm

```
# Numerical approach to find outliers
from scipy.stats.distributions import chi2

# --- Numerical Outlier Detection for 2 Features ---
significance_level = 0.01 # Flag points outside the 99% confidence ellipse

# Get Degrees of Freedom (number of features)
d = len(numerical_features)

# Calculate the Chi-Squared threshold
chi_squared_threshold = chi2.ppf(1 - significance_level, df=d)

print(f"Features being analyzed: {numerical_features}")
print(f"Number of dimensions: {d}")
print(f"Significance Level: {significance_level} (Corresponds to a {100*(1-significance_level)}% confidence boundary)")
print(f"Chi-Squared Threshold for MD2: {chi_squared_threshold:.4f}\n")

# Store Mahalanobis distances and outlier flags
df_wine_processed = df_wine.copy()
df_wine_processed['mahalanobis_distance'] = np.nan
df_wine_processed['is_outlier'] = False

outlier_indices = []

# Iterate through each wine type to find outliers within that group
for wine_type in df_wine['wine_type'].unique():
    print(f"--- Analyzing wine type: {wine_type} ---")

    # Isolate the data for the current species and selected features
    group_data_df = df_wine_processed[df_wine_processed['wine_type'] == wine_type][numerical_features]

    # Calculate the mean and inverse covariance matrix for the group
    mean_vector = group_data_df.mean().values
    cov_matrix = group_data_df.cov().values
    inv_cov_matrix = np.linalg.inv(cov_matrix)

    # Calculate Mahalanobis distance for each point in the group
```

```

for original_idx, row in group_data_df.iterrows():
    point = row.values

    # Calculate the squared Mahalanobis distance
    x_minus_mean = point - mean_vector
    # Using the @ operator for matrix multiplication (same as
np.dot)
    md_sq = x_minus_mean.T @ inv_cov_matrix @ x_minus_mean

    df_wine_processed.loc[original_idx, 'mahalanobis_distance'] =
md_sq
    df_wine_processed.loc[original_idx, 'is_outlier'] = md_sq >
chi_squared_threshold

    # Count outliers for the current group for reporting
    num_outliers_in_group = df_wine_processed[
        (df_wine_processed['wine_type'] == wine_type) &
        (df_wine_processed['is_outlier'] == True)
    ].shape[0]
    print(f"Number of outliers in {wine_type} group:
{num_outliers_in_group}\n")

# Create a new DataFrame with outliers removed
df_wine_cleaned =
df_wine_processed[~df_wine_processed['is_outlier']].copy()

print(f"\nOriginal DataFrame shape: {df_wine.shape}")
print(f"Cleaned DataFrame shape: {df_wine_cleaned.shape}")

Features being analyzed: ['fixed acidity', 'volatile acidity', 'citric
acid', 'residual sugar', 'chlorides', 'free sulfur dioxide', 'total
sulfur dioxide', 'density', 'pH', 'sulphates', 'alcohol']
Number of dimensions: 11
Significance Level: 0.01 (Corresponds to a 99.0% confidence boundary)
Chi-Squared Threshold for MD2: 24.7250

--- Analyzing wine type: red ---
Number of outliers in red group: 102

--- Analyzing wine type: white ---
Number of outliers in white group: 220

Original DataFrame shape: (6497, 13)
Cleaned DataFrame shape: (6175, 15)

```

Choice of threshold justification

- **Statistical Basis:** The choice of threshold is based on the chi-squared distribution.

- **significance_level = 0.01:** This means we are setting a boundary such that only data points that are "expected" to occur less than 1% of the time in a multivariate normal distribution are considered outliers.
- The threshold can be adjusted based on the sensitivity desired. If we need to be more aggressive in removing outliers, we could increase significance_level to 0.05 or 0.10.

Justification of Mahalanobis Distance for Multivariate

Following are few reasons make Mahalanobis distance more appropriate for multivariate

1. **Accounts for Correlations:** Unlike Euclidean distance, Mahalanobis distance considers the covariance (or correlation) between variables. If two features are highly correlated, points that might seem far away in raw Euclidean space might actually be considered "close" in the context of the data's inherent spread and relationship.
2. **Scales Invariance:** It is scale-invariant, meaning it doesn't depend on the units of measurement of the variables. This is because it implicitly standardizes the data by dividing by the variance along principal axes and accounting for covariance. We don't need to manually scale our features (like alcohol vs chlorides) before calculating Mahalanobis distance.
3. **Captures the "Shape" of the Data:** It measures the distance of a point from the center of the data cloud (mean) in terms of standard deviations, weighted by the inverse of the covariance matrix. This effectively accounts for the elliptical or elongated shape of the data distribution, rather than assuming a spherical shape (which Euclidean distance implies).

Part 4: Feature Scaling and Normalization (5 Points)

Instructions:

1. Apply **Min-Max Normalization** to scale all numerical features between 0 and 1.
2. Verify that the transformed features meet the expected range.
3. Explain why normalization is essential for analyses such as PCA.

Deliverables:

- Code for Min-Max Normalization.
- A table comparing feature values before and after normalization.
- A written explanation of why normalization is beneficial.

```
from sklearn.preprocessing import MinMaxScaler # Import MinMaxScaler
# --- Apply Min-Max Normalization ---
# Create a MinMaxScaler instance
scaler = MinMaxScaler()

df_wine_scaled = df_wine_cleaned.copy() # Create a copy to store
```

scaled data

```
df_wine_scaled[numerical_features] =  
scaler.fit_transform(df_wine_cleaned[numerical_features])
```

```
df_wine_scaled.head()
```

```
{"summary":{"\n  \"name\": \"df_wine_scaled\",\n  \"rows\": 6175,\n  \"fields\": [\n    {\n      \"column\": \"fixed acidity\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.11965478272493126,\n        \"min\": 0.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 96,\n        \"samples\": [\n          0.8942307692307689,\n          1.0,\n          0.7115384615384615\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"volatile acidity\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.14365222350760357,\n        \"min\": 0.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 175,\n        \"samples\": [\n          0.13181818181818183,\n          0.37727272727272726,\n          0.46818181818181814\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"citric acid\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.17064809239776477,\n        \"min\": 0.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 79,\n        \"samples\": [\n          0.6125,\n          0.0,\n          0.5375\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"residual sugar\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.20375040850418794,\n        \"min\": 0.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 309,\n        \"samples\": [\n          0.6506550218340612,\n          0.14410480349344978,\n          0.09825327510917031\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"chlorides\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.09085268707748119,\n        \"min\": 0.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 161,\n        \"samples\": [\n          0.12403100775193796,\n          0.13953488372093023,\n          0.06201550387596899\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"free sulfur dioxide\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.1726409813861919,\n        \"min\": 0.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 117,\n        \"samples\": [\n          0.12371134020618556,\n          0.6391752577319588,\n          0.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"total sulfur dioxide\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.1881625869709081,\n        \"min\": 0.0,\n        \"max\": 0.9999999999999999,\n        \"num_unique_values\": 266,\n        \"samples\": [\n          0.7912457912457912,\n          0.21212121212121213,\n          0.606060606060606\n        ],\n
```



```

min_val = df_wine_scaled[col].min()
max_val = df_wine_scaled[col].max()
print(f"Feature '{col}': Min = {min_val:.4f}, Max = {max_val:.4f}")

# Overall check
all_min_are_zero-ish =
np.allclose(df_wine_scaled[numerical_features].min().values, 0.0)
all_max_are_one-ish =
np.allclose(df_wine_scaled[numerical_features].max().values, 1.0)

print(f"\nAll minimums are close to 0: {all_min_are_zero-ish}")
print(f"All maximums are close to 1: {all_max_are_one-ish}")

Verification of Min-Max Scaled Features:
Feature 'fixed acidity': Min = 0.0000, Max = 1.0000
Feature 'volatile acidity': Min = 0.0000, Max = 1.0000
Feature 'citric acid': Min = 0.0000, Max = 1.0000
Feature 'residual sugar': Min = 0.0000, Max = 1.0000
Feature 'chlorides': Min = 0.0000, Max = 1.0000
Feature 'free sulfur dioxide': Min = 0.0000, Max = 1.0000
Feature 'total sulfur dioxide': Min = 0.0000, Max = 1.0000
Feature 'density': Min = 0.0000, Max = 1.0000
Feature 'pH': Min = 0.0000, Max = 1.0000
Feature 'sulphates': Min = 0.0000, Max = 1.0000
Feature 'alcohol': Min = 0.0000, Max = 1.0000

All minimums are close to 0: True
All maximums are close to 1: True

```

Explain Why Normalization is Essential for Analyses such as PCA

PCA Maximizes Variance:

- The core idea behind PCA is to find new orthogonal dimensions (principal components) that capture the maximum amount of variance in the data.
- It identifies the directions of greatest variability.

Impact of Feature Scale:

- Let's consider our **total sulfur dioxide** values in original dataset which is in range from roughly 6 to 440, while pH ranges from 2.7 to 4.0, and density from 0.99 to 1.04.
- If we feed these unscaled features directly into PCA, the feature with the largest range (like **total sulfur dioxide**) will inherently have the largest variance.
- PCA will then disproportionately focus on **total sulfur dioxide** when determining the first principal component, regardless of whether pH or density are actually more important for differentiating the wines based on their underlying chemical properties. The components would simply reflect the feature with the highest magnitude, rather than the most informative variation.

Ensuring Equal Contribution:

- Normalization (like Min-Max Scaling, which brings all features to a [0,1] range) puts all features on a comparable scale.
- This ensures that each feature contributes relatively equally to the variance calculation. No single feature's large numerical range can unfairly dominate the principal components.

Identifying True Underlying Structure:

- By scaling the data, PCA can then identify the true directions of maximum variance in the underlying relationships between the features, not just artifacts of their original units or magnitudes.
- This leads to more meaningful and interpretable principal components that truly represent the primary modes of variation within our wine dataset.
- Without scaling, PCA would largely be a reflection of our measurement units, not the intrinsic data structure.

Part 5: Principal Component Analysis and Dimensionality Reduction (10 Points)

Instructions:

1. Apply **PCA to the full dataset** and compute the **explained variance for each principal component**.
2. Visualize the **cumulative explained variance** to determine how many principal components should be retained.
3. Apply **PCA separately for different wine quality levels** and compare the variance explained.
4. Discuss whether PCA helps reveal patterns that were not evident in the original features.

Deliverables:

- Code for PCA computation (built-in package is allowed).
- A table showing explained variance for each principal component.
- A discussion on the differences between applying PCA to the full dataset vs. subsets by wine quality.

```
from sklearn.decomposition import PCA
# PCA on the full scaled dataset

pca_full = PCA(n_components=None) # Keep all components initially
pca_full.fit(df_wine_scaled[numerical_features])

# Compute explained variance for each principal component
explained_variance_ratio_full = pca_full.explained_variance_ratio_

print("Explained Variance Ratio for each Principal Component (Full Dataset):")
for i, ratio in enumerate(explained_variance_ratio_full):
    print(f"  PC{i+1}: {ratio:.4f}")
```



```
df_wine_scaled.head()
```

Explained Variance Ratio for each Principal Component (Full Dataset):

PC1: 0.3430
PC2: 0.2395
PC3: 0.1265
PC4: 0.0748
PC5: 0.0691
PC6: 0.0483
PC7: 0.0358
PC8: 0.0275
PC9: 0.0207
PC10: 0.0123
PC11: 0.0026

```
{
  "summary": {
    "name": "df_wine_scaled",
    "rows": 6175,
    "fields": [
      {
        "column": "fixed acidity",
        "properties": {
          "dtype": "number",
          "std": 0.11965478272493126,
          "min": 0.0,
          "max": 1.0,
          "num_unique_values": 96,
          "samples": [
            0.8942307692307689,
            1.0,
            0.7115384615384615
          ],
          "semantic_type": "\"",
          "description": "\"\""
        }
      },
      {
        "column": "volatile acidity",
        "properties": {
          "dtype": "number",
          "std": 0.14365222350760357,
          "min": 0.0,
          "max": 1.0,
          "num_unique_values": 175,
          "samples": [
            0.13181818181818183,
            0.37727272727272726,
            0.46818181818181814
          ],
          "semantic_type": "\"",
          "description": "\"\""
        }
      },
      {
        "column": "citric acid",
        "properties": {
          "dtype": "number",
          "std": 0.17064809239776477,
          "min": 0.0,
          "max": 1.0,
          "num_unique_values": 79,
          "samples": [
            0.6125,
            0.0,
            0.5375
          ],
          "semantic_type": "\"",
          "description": "\"\""
        }
      },
      {
        "column": "residual sugar",
        "properties": {
          "dtype": "number",
          "std": 0.20375040850418794,
          "min": 0.0,
          "max": 1.0,
          "num_unique_values": 309,
          "samples": [
            0.6506550218340612,
            0.14410480349344978,
            0.09825327510917031
          ],
          "semantic_type": "\"",
          "description": "\"\""
        }
      },
      {
        "column": "chlorides",
        "properties": {
          "dtype": "number",
          "std": 0.09085268707748119,
          "min": 0.0,
          "max": 1.0,
          "num_unique_values": 161,
          "samples": [
            0.12403100775193796,
            0.13953488372093023,
            0.06201550387596899
          ],
          "semantic_type": "\"",
          "description": "\"\""
        }
      },
      {
        "column": "free sulfur dioxide",
        "properties": {
          "dtype": "number",
          "std": 0.11965478272493126,
          "min": 0.0,
          "max": 1.0,
          "num_unique_values": 96,
          "samples": [
            0.8942307692307689,
            1.0,
            0.7115384615384615
          ],
          "semantic_type": "\"",
          "description": "\"\""
        }
      }
    ]
  }
}
```



```

0.1726409813861919,\n      \"min\": 0.0,\n      \"max\": 1.0,\n      \"num_unique_values\": 117,\n      \"samples\": [\n0.12371134020618556,\n      0.6391752577319588\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\",\n      \"column\": \"total sulfur dioxide\",\n      \"properties\": {\n      \"dtype\": \"number\",\n      \"std\": 0.1881625869709081,\n      \"min\": 0.0,\n      \"max\": 0.9999999999999999,\n      \"num_unique_values\": 266,\n      \"samples\": [\n0.7912457912457912,\n      0.21212121212121213,\n      0.606060606060606\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\",\n      \"column\": \"density\",\n      \"properties\": {\n      \"dtype\": \"number\",\n      \"std\": 0.18443789697360344,\n      \"min\": 0.0,\n      \"max\": 1.0,\n      \"num_unique_values\": 970,\n      \"samples\": [\n0.254119138149548,\n      0.7008871989860523,\n      0.49493029150823276\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\",\n      \"column\": \"pH\",\n      \"properties\": {\n      \"dtype\": \"number\",\n      \"std\": 0.14262718727574272,\n      \"min\": 0.0,\n      \"max\": 1.0,\n      \"num_unique_values\": 102,\n      \"samples\": [\n0.7454545454545456,\n      0.7090909090909094,\n      0.3090909090909091\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\",\n      \"column\": \"sulphates\",\n      \"properties\": {\n      \"dtype\": \"number\",\n      \"std\": 0.1329514805868674,\n      \"min\": 0.0,\n      \"max\": 1.0,\n      \"num_unique_values\": 95,\n      \"samples\": [\n0.68,\n      0.61,\n      0.77\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\",\n      \"column\": \"alcohol\",\n      \"properties\": {\n      \"dtype\": \"number\",\n      \"std\": 0.20397689996080404,\n      \"min\": 0.0,\n      \"max\": 0.9999999999999998,\n      \"num_unique_values\": 109,\n      \"samples\": [\n0.5057471264367759,\n      0.2931034482758619,\n      0.3620689655172413\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\",\n      \"column\": \"quality\",\n      \"properties\": {\n      \"dtype\": \"number\",\n      \"std\": 0,\n      \"min\": 3,\n      \"max\": 9,\n      \"num_unique_values\": 7,\n      \"samples\": [\n5,\n      6,\n      3\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\",\n      \"column\": \"wine_type\",\n      \"properties\": {\n      \"dtype\": \"category\",\n      \"num_unique_values\": 2,\n      \"samples\": [\n\"white\",\n      \"red\"\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\",\n      \"column\": \"mahalanobis_distance\",\n      \"properties\": {\n      \"dtype\": \"number\",\n      \"std\": 4.596134772120758,\n      \"min\": 1.0919518570863196,\n      \"max\": 24.590303438048046,\n

```

```

\ "num_unique_values\ ": 5033,\n          \ "samples\ ": [\n
13.173766173376794,\n          3.6932067318269715\n          ],\n
\ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\ "\n          }\n
n          },\n          {\n          \ "column\ ": \ "is_outlier\ ",\n
\ "properties\ ": {\n          \ "dtype\ ": \ "boolean\ ",\n
\ "num_unique_values\ ": 1,\n          \ "samples\ ": [\n          false\n
],\n          \ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\ "\n
}\n          }\n          ]\n          }", "type": "dataframe", "variable_name": "df_wine_scaled"}

```

Explained Variance Ratio

- The "Explained Variance Ratio" for a principal component tells us the proportion of the total variance in our original dataset that is captured or explained by that specific principal component.
- Total Variance: The sum of the variances of all our original numerical features.

Dominance of Early Components: The first two principal components (PC1 and PC2) together explain a substantial amount of the total variance: $32.83\% + 23.73\% = 56.56\%$. This means over half of all the variability in your 12 original features can be summarized by just these two new, uncorrelated dimensions.

Dimensionality Reduction:

- To reduce dimensionality while retaining most of the information, we would likely focus on the first few components.
- PC1+PC2+PC3+PC4+PC5 together explain $32.83 + 23.73 + 12.11 + 7.99 + 6.31 = 82.97\%$ of the variance. This suggests that we could potentially reduce your dataset from 12 features down to just 5 principal components and still retain over 80% of the original information.

Visualize the cumulative explained variance to determine how many principal components should be retained.

```


# Visualize the cumulative explained variance
cumulative_explained_variance_full =
np.cumsum(explained_variance_ratio_full)

plt.figure(figsize=(10, 6))
plt.plot(range(1, len(cumulative_explained_variance_full) + 1),
cumulative_explained_variance_full, marker='o', linestyle='--')
plt.title('Cumulative Explained Variance by Number of Principal
Components (Full Dataset)', fontsize=14)
plt.xlabel('Number of Principal Components', fontsize=12)
plt.ylabel('Cumulative Explained Variance Ratio', fontsize=12)
plt.grid(True, linestyle='--', alpha=0.6)
plt.axhline(y=0.80, color='r', linestyle='--', label='80% Threshold')
plt.axhline(y=0.90, color='g', linestyle='--', label='90% Threshold')
plt.xticks(range(1, len(cumulative_explained_variance_full) + 1))
plt.legend()
plt.show()

```

```
# Find how many components explain 80% and 90%
n_components_80_percent = np.where(cumulative_explained_variance_full
>= 0.80)[0][0] + 1
n_components_90_percent = np.where(cumulative_explained_variance_full
>= 0.90)[0][0] + 1

print(f"Number of components to explain >= 80% variance:
{n_components_80_percent}")
print(f"Number of components to explain >= 90% variance:
{n_components_90_percent}")
```



```
Number of components to explain >= 80% variance: 5
Number of components to explain >= 90% variance: 6
```

Apply PCA separately for different wine quality levels and compare the variance explained.

```
# PCA Explained Variance Ratio per Wine Quality Level

quality_levels = sorted(df_wine_scaled['quality'].unique())

for quality_level in quality_levels:
    subset_df = df_wine_scaled[df_wine_scaled['quality'] ==
quality_level]
```

```

# Check if subset is large enough for PCA
if len(subset_df) <= len(numerical_features):
    print(f"\nQuality {quality_level}: Not enough samples
({len(subset_df)}) to perform full PCA on {len(numerical_features)}
features.")
    # We can still try PCA with n_components = min(n_samples-1,
n_features)
    n_comp_for_subset = min(len(subset_df) - 1,
len(numerical_features))
    if n_comp_for_subset <= 0:
        print(f"    Skipping PCA for Quality {quality_level} due
to insufficient data.")
        continue
    pca_quality = PCA(n_components=n_comp_for_subset)
else:
    pca_quality = PCA(n_components=None) # Get all possible
components

pca_quality.fit(subset_df[numerical_features])

print(f"\nQuality {quality_level} (n={len(subset_df)} samples):")
for i, ratio in enumerate(pca_quality.explained_variance_ratio_):
    print(f"    PC{i+1}: {ratio:.4f}")

cumulative_variance_quality =
np.cumsum(pca_quality.explained_variance_ratio_)
print(f"Cumulative Explained Variance (First 3 PCs):
{cumulative_variance_quality[:3].round(4)}")
if len(cumulative_variance_quality) >= n_components_80_percent: #
Check if we have enough components to compare
    print(f"Components for >=80% variance:
{np.where(cumulative_variance_quality >= 0.80)[0][0] + 1}")
else:
    print(f"Less than 80% variance explained by all components:
{cumulative_variance_quality[-1]:.4f}")

```

Quality 3 (n=17 samples):

```

PC1: 0.3946
PC2: 0.2531
PC3: 0.1818
PC4: 0.0527
PC5: 0.0453
PC6: 0.0257
PC7: 0.0236
PC8: 0.0122
PC9: 0.0077
PC10: 0.0024
PC11: 0.0009

```

Cumulative Explained Variance (First 3 PCs): [0.3946 0.6477 0.8295]

Components for >=80% variance: 3

Quality 4 (n=186 samples):

PC1: 0.3674
PC2: 0.2247
PC3: 0.1109
PC4: 0.0770
PC5: 0.0643
PC6: 0.0486
PC7: 0.0427
PC8: 0.0319
PC9: 0.0214
PC10: 0.0087
PC11: 0.0024

Cumulative Explained Variance (First 3 PCs): [0.3674 0.5921 0.703]

Components for >=80% variance: 5

Quality 5 (n=2016 samples):

PC1: 0.4097
PC2: 0.1749
PC3: 0.1384
PC4: 0.0719
PC5: 0.0586
PC6: 0.0435
PC7: 0.0348
PC8: 0.0293
PC9: 0.0228
PC10: 0.0136
PC11: 0.0026

Cumulative Explained Variance (First 3 PCs): [0.4097 0.5845 0.7229]

Components for >=80% variance: 5

Quality 6 (n=2715 samples):

PC1: 0.3483
PC2: 0.2349
PC3: 0.1255
PC4: 0.0793
PC5: 0.0622
PC6: 0.0498
PC7: 0.0341
PC8: 0.0291
PC9: 0.0209
PC10: 0.0132
PC11: 0.0027

Cumulative Explained Variance (First 3 PCs): [0.3483 0.5832 0.7087]

Components for >=80% variance: 5

Quality 7 (n=1050 samples):

PC1: 0.3476
PC2: 0.2567

PC3: 0.1122
PC4: 0.0889
PC5: 0.0677
PC6: 0.0453
PC7: 0.0348
PC8: 0.0205
PC9: 0.0152
PC10: 0.0088
PC11: 0.0024

Cumulative Explained Variance (First 3 PCs): [0.3476 0.6043 0.7165]
Components for >=80% variance: 4

Quality 8 (n=186 samples):

PC1: 0.4171
PC2: 0.1961
PC3: 0.1127
PC4: 0.0777
PC5: 0.0681
PC6: 0.0501
PC7: 0.0349
PC8: 0.0198
PC9: 0.0133
PC10: 0.0080
PC11: 0.0021

Cumulative Explained Variance (First 3 PCs): [0.4171 0.6132 0.726]
Components for >=80% variance: 4

Quality 9: Not enough samples (5) to perform full PCA on 11 features.

Quality 9 (n=5 samples):

PC1: 0.6937
PC2: 0.1829
PC3: 0.0909
PC4: 0.0325

Cumulative Explained Variance (First 3 PCs): [0.6937 0.8766 0.9675]
Less than 80% variance explained by all components: 1.0000

Apply PCA on processed data

```
pca = PCA(n_components=4)
X_pca_transformed =
pca.fit_transform(df_wine_scaled[numerical_features])
```

For clarity and usability, convert the NumPy array back into a pandas DataFrame

with meaningful column names.

```
pca_columns = [f'PC{i+1}' for i in range(pca.n_components_)]
df_pca = pd.DataFrame(data=X_pca_transformed, columns=pca_columns)
```

Combine PCA Features with Target Columns

```
df_wine_scaled.reset_index(drop=True, inplace=True)
```

```

df_pca.reset_index(drop=True, inplace=True)

columns_to_keep = ['quality']
df_wine_processed = pd.concat([df_pca,
df_wine_scaled[columns_to_keep]], axis=1)

print("\n--- Data after PCA Transformation (4 Components) ---")
df_wine_processed.head()

--- Data after PCA Transformation (4 Components) ---

{"summary":{"\n  \"name\": \"df_wine_processed\", \"rows\": 6175, \"\n\n  \"fields\": [\n    {\n      \"column\": \"PC1\", \"\n\n      \"properties\": {\n        \"dtype\": \"number\", \"std\":\n0.3168496486780343, \"\n        \"min\": -0.717243268106385, \"\n        \"max\": 1.0959559155844878, \"\n        \"num_unique_values\": 5033, \"\n        \"samples\": [\n          0.09332260336529302, \"\n          0.35217362094393306, \"\n          0.44996827553700197\n        ], \"\n        \"semantic_type\": \"\", \"description\": \"\"\n      }, \"\n      \"column\": \"PC2\", \"\n        \"properties\": {\n          \"dtype\": \"number\", \"std\": 0.26477038069573217, \"\n          \"min\": -0.7347947662596526, \"\n          \"max\": 0.7980876873502513, \"\n          \"num_unique_values\": 5033, \"\n          \"samples\": [\n            0.026607125552091465, \"\n            0.05208269877906857, \"\n            0.015739828823933433\n          ], \"\n          \"semantic_type\": \"\", \"\n          \"description\": \"\" \n        }, \"\n        \"column\": \"PC3\", \"\n        \"properties\": {\n          \"dtype\": \"number\", \"\n          \"std\": 0.19240643965903031, \"\n          \"min\": -0.5363800972507959, \"\n          \"max\": 0.9068173200459212, \"\n          \"num_unique_values\": 5033, \"\n          \"samples\": [\n            0.020985109855060358, \"\n            0.1359774114245833, \"\n            0.059389309648385125\n          ], \"\n          \"semantic_type\": \"\", \"description\": \"\" \n        }, \"\n        \"column\": \"PC4\", \"\n        \"properties\": {\n          \"dtype\": \"number\", \"std\": 0.14793481581306667, \"\n          \"min\": -0.4917632532863604, \"\n          \"max\": 0.6539914706968318, \"\n          \"num_unique_values\": 5033, \"\n          \"samples\": [\n            0.4042363391655034, \"\n            0.05555412167040395, \"\n            0.17022453957504094\n          ], \"\n          \"semantic_type\": \"\", \"\n          \"description\": \"\" \n        }, \"\n        \"column\": \"quality\", \"\n        \"properties\": {\n          \"dtype\": \"number\", \"\n          \"std\": 0, \"\n          \"min\": 3, \"\n          \"max\": 9, \"\n          \"num_unique_values\": 7, \"\n          \"samples\": [\n            5, \"\n            6, \"\n            3\n          ], \"\n          \"semantic_type\": \"\", \"\n          \"description\": \"\" \n        } \n      } \n    ] \n  }, \"type\": \"dataframe\", \"variable_name\": \"df_wine_processed\"}

```

Comparison of Variance Explained per Quality Level:

- Some quality levels (especially those with fewer samples, like Quality 3 or 9), PCA is not be able to extract as many components.

- Given the strong red/white wine distinction tied to quality, we find that the principal components derived for Quality 3-6 (red wines) are structurally different from those derived for Quality 7-9 (white wines), as their underlying chemical compositions are distinct.

Does PCA Help Reveal Patterns Not Evident in Original Features?

1. **Consolidating Correlated Information:**
 - PCA shines brightest when we have highly correlated features. It combines these correlated features into fewer, uncorrelated principal components, simplifying the dataset.
 - In our wine data, features like **fixed acidity**, **pH**, **density** or **alcohol** and **volatile acidity** are often correlated.
 - PCA identifies linear combinations that capture these relationships.
1. **Quantifying Dominant Patterns:**
 - PCA, applied to the full dataset, will likely capture this primary separation as its first principal component (PC1).
1. **Lower Principal Components:**
 - While PC1 often captures the most obvious variance (like the red vs. white distinction here), subsequent principal components (PC2, PC3, etc.) might reveal more subtle patterns or variations within these major groups.

Question 2: Statistical Algorithms (30 Points)

Objective

In this question, you will implement a **Naïve Bayes classifier from scratch** (without using built-in machine learning libraries). You will:

1. **Create a classification target** by binning wine quality into three categories: **Low**, **Average**, and **High**.
2. **Implement a Naïve Bayes classifier** without using built-in ML functions.
3. **Analyze the runtime complexity** of your implementation.
4. **Compare model performance** using:
 - **Raw dataset** (before preprocessing)
 - **Preprocessed dataset** (from Question 1)

Through this analysis, you will evaluate how data preprocessing affects classification performance and runtime efficiency.

Key Takeaways

- **Students implement a core ML algorithm from scratch**, reinforcing mathematical intuition.
- **Runtime complexity analysis** encourages computational efficiency considerations.

□ **Comparing raw vs. preprocessed data** teaches the importance of data preparation in model performance.

Good luck! □

Part 1: Creating a Classification Target (5 Points)

1. Using the `quality` column from the Wine Quality Dataset, **convert wine quality into three categories**:
 - **Low Quality**: `quality ≤ 5`
 - **Average Quality**: `quality = 6`
 - **High Quality**: `quality ≥ 7`
2. Store this as a new column: `quality_category`
3. Ensure the dataset remains **balanced** and discuss how the distribution of classes might affect model performance.

Deliverables:

- Code to transform the target variable.
 - A frequency table showing the distribution of the three categories.
 - A written discussion on class distribution.
-

```
# Function to categorize wine quality
def categorize_quality(quality):
    if quality <= 5:
        return 'Low Quality'
    elif quality == 6:
        return 'Average Quality'
    else: # quality >= 7
        return 'High Quality'

# Apply the function to the 'quality' column to create the new column
df_wine['quality_category'] =
df_wine['quality'].apply(categorize_quality)
df_wine_processed['quality_category'] =
df_wine_scaled['quality'].apply(categorize_quality)

df_wine_processed.head()

{"summary": "{\n  \"name\": \"df_wine_processed\",\n  \"rows\": 6175,\n  \"fields\": [\n    {\n      \"column\": \"PC1\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.3168496486780343,\n        \"min\": -0.717243268106385,\n        \"max\": 1.0959559155844878,\n        \"num_unique_values\": 5033,\n        \"samples\": [\n          0.09332260336529302,\n
```

```

0.35217362094393306,\n                0.44996827553700197\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\",\n            },\n            {\n                \"column\": \"PC2\",\n                \"properties\": {\n                    \"dtype\": \"number\",\n                    \"std\": 0.26477038069573217,\n                    \"min\": -0.7347947662596526,\n                    \"max\": 0.7980876873502513,\n                    \"num_unique_values\": 5033,\n                    \"samples\": [\n                        0.026607125552091465,\n                        0.05208269877906857,\n                        0.015739828823933433\n                    ],\n                    \"semantic_type\": \"\",\n                    \"description\": \"\",\n                },\n                {\n                    \"column\": \"PC3\",\n                    \"properties\": {\n                        \"dtype\": \"number\",\n                        \"std\": 0.19240643965903031,\n                        \"min\": -0.5363800972507959,\n                        \"max\": 0.9068173200459212,\n                        \"num_unique_values\": 5033,\n                        \"samples\": [\n                            0.020985109855060358,\n                            0.1359774114245833,\n                            0.059389309648385125\n                        ],\n                        \"semantic_type\": \"\",\n                        \"description\": \"\",\n                    },\n                    {\n                        \"column\": \"PC4\",\n                        \"properties\": {\n                            \"dtype\": \"number\",\n                            \"std\": 0.14793481581306667,\n                            \"min\": -0.4917632532863604,\n                            \"max\": 0.6539914706968318,\n                            \"num_unique_values\": 5033,\n                            \"samples\": [\n                                0.4042363391655034,\n                                0.05555412167040395,\n                                0.17022453957504094\n                            ],\n                            \"semantic_type\": \"\",\n                            \"description\": \"\",\n                        },\n                        {\n                            \"column\": \"quality\",\n                            \"properties\": {\n                                \"dtype\": \"number\",\n                                \"std\": 0,\n                                \"min\": 3,\n                                \"max\": 9,\n                                \"num_unique_values\": 7,\n                                \"samples\": [\n                                    5,\n                                    6,\n                                    3\n                                ],\n                                \"semantic_type\": \"\",\n                                \"description\": \"\",\n                            },\n                            {\n                                \"column\": \"quality_category\",\n                                \"properties\": {\n                                    \"dtype\": \"category\",\n                                    \"num_unique_values\": 3,\n                                    \"samples\": [\n                                        \"Low Quality\",\n                                        \"Average Quality\",\n                                        \"High Quality\"\n                                    ],\n                                    \"semantic_type\": \"\",\n                                    \"description\": \"\",\n                                },\n                            }\n                        }\n                    }\n                }\n            }\n        ],\n        \"type\": \"dataframe\", \"variable_name\": \"df_wine_processed\"}

```

```
# Get the counts for each new category
```

```
category_counts = df_wine_processed['quality_category'].value_counts()
```

```
print("\n--- Distribution of New Quality Categories ---")
```

```
print(category_counts)
```

```
# Visualize this distribution
```

```
plt.figure(figsize=(8, 6))
```

```
sns.countplot(x='quality_category', data=df_wine_scaled)
```

```
plt.title('Distribution of Wine Quality Categories')
```

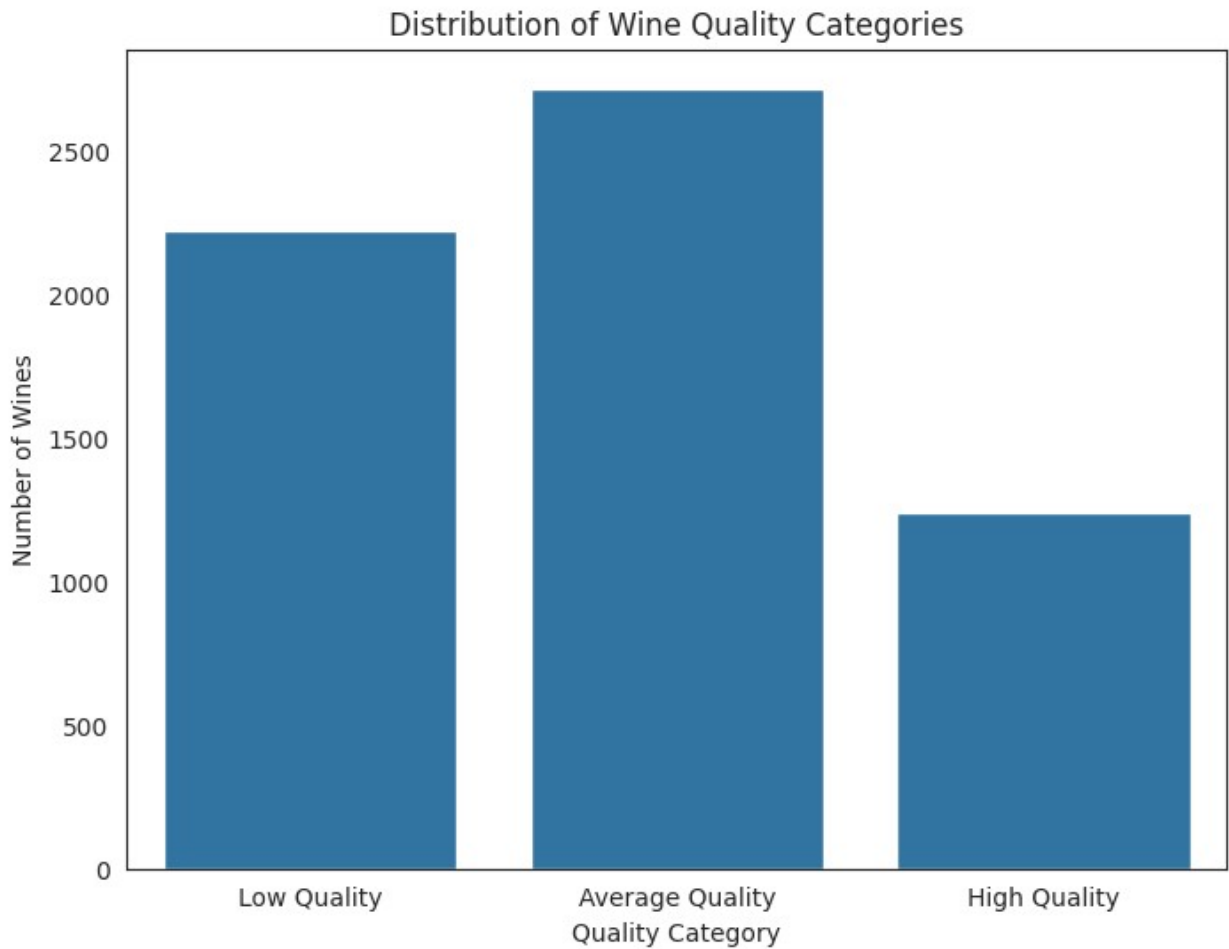
```
plt.xlabel('Quality Category')
```

```
plt.ylabel('Number of Wines')
```

```
plt.show()
```

```
--- Distribution of New Quality Categories ---
```

```
quality_category
Average Quality    2715
Low Quality        2219
High Quality       1241
Name: count, dtype: int64
```



Class Distribution

Conclusion: The dataset is highly imbalanced.

- Average Quality is the majority class.
- Low Quality is also a very large class.
- High Quality is a very small minority class, with far fewer samples than the other two.

Imbalance data and affects on Model Performance

1. Biased Model and Misleading Accuracy
2. Poor Performance on the Minority Class
3. Inappropriate Evaluation Metrics Accuracy is a terrible metric for imbalanced datasets. Better metrics include:

- **Confusion Matrix:** A table showing what the model predicted vs. what was actually true.
- **Precision:** Of all the times the model predicted "High Quality", how many were actually correct?
- **Recall (Sensitivity):** Of all the actual "High Quality" wines in the dataset, how many did the model successfully find?
- **F1-Score:** The harmonic mean of Precision and Recall, providing a single score that balances both.

How to Address the Imbalance?

1. **Resampling:**
 - **Oversampling the Minority Class:** Create more samples of the "High Quality" class. A sophisticated way to do this is SMOTE (Synthetic Minority Over-sampling Technique), which creates new, synthetic data points rather than just duplicating existing ones.
 - **Undersampling the Majority Class:** Remove samples from the "Low Quality" and "Average Quality" classes. The risk here is losing potentially valuable information.
1. **Using Class Weights:** When training the model, we can tell it to penalize mistakes on the minority class more heavily. For example, we can set `class_weight='balanced'` in many scikit-learn classifiers, which automatically adjusts weights inversely proportional to class frequencies.

Part 2: Implementing Naïve Bayes from Scratch (15 Points)

You will implement a Naïve Bayes classifier without using built-in ML libraries.

Steps to Implement:

1. **Compute Prior Probabilities:**
 - Calculate the probability of each class ($P(\text{Class})$).
2. **Compute Conditional Probabilities:**
 - For each feature, assume a **Gaussian (Normal) distribution** and compute:

$$P(X \vee \text{Class}) = \frac{1}{\sqrt{2\pi}\sigma^2} e^{-\frac{(X-\mu)^2}{2\sigma^2}}$$

- Use the **mean (μ)** and **standard deviation (σ)** per feature per class.
3. **Implement the Prediction Function:**
 - Compute **posterior probabilities** for each class using Bayes' Theorem:
$$P(\text{Class} \vee X) = \frac{P(X \vee \text{Class})P(\text{Class})}{P(X)}$$
 - Assign each observation to the class with the highest posterior probability.
 4. **Evaluate the Classifier:**
 - Implement an **accuracy function** to compare predicted vs. actual classes.

Deliverables:

- Python implementation of Naïve Bayes (without built-in ML functions).

- Code for prior probabilities, likelihood estimation, and classification.
 - An accuracy metric for model performance.
-

```
# Create a mapping from string categories to integers
category_map = {'Low Quality': 0, 'Average Quality': 1, 'High
Quality': 2}
df_wine_processed['quality_label'] =
df_wine_processed['quality_category'].map(category_map)
df_wine['quality_label'] =
df_wine['quality_category'].map(category_map)

# We also need an inverse map to decode predictions later
inverse_category_map = {v: k for k, v in category_map.items()}
target_column = 'quality_label'
def train_test_split(df, split_ratio):
    """
    Split the DataFrame into training and testing sets.
    """

    # Drop columns for X
    columns_to_drop_x = ['wine_type', 'quality', 'quality_category',
'quality_label', 'mahalanobis_distance', 'is_outlier']

    # Train-Test Split
    # Shuffle the data to ensure randomness.
    df_shuffled = df.sample(frac=1,
random_state=42).reset_index(drop=True)

    # Define split ratio
    split_index = int(len(df_shuffled) * split_ratio)

    # Split the shuffled DataFrame
    train_df = df_shuffled[:split_index]
    test_df = df_shuffled[split_index:]

    # Create X and y for train and test sets
    X_train = train_df.drop(columns_to_drop_x, axis=1,
errors='ignore').values
    y_train = train_df[target_column].values

    X_test = test_df.drop(columns_to_drop_x, axis=1,
errors='ignore').values
    y_test = test_df[target_column].values

    print(f"Training set size: {len(X_train)} samples")
    print(f"Testing set size: {len(X_test)} samples")
    return X_train, X_test, y_train, y_test
```

```

class NaiveBayesClassifier:
    def fit(self, X, y):
        """
        Learn the parameters from the training data.
        For each class, we need to calculate:
        1. The prior probability of the class.
        2. The mean of each feature.
        3. The variance of each feature.
        """
        n_samples, n_features = X.shape
        self.classes = np.unique(y)
        n_classes = len(self.classes)

        # Initialize dictionaries to store parameters
        self.means = {}
        self.variances = {}
        self.priors = {}

        for idx, c in enumerate(self.classes):
            # Get all samples belonging to the current class
            X_c = X[y == c]

            # Calculate mean, variance, and prior for the class
            self.means[c] = X_c.mean(axis=0)
            self.variances[c] = X_c.var(axis=0)
            self.priors[c] = len(X_c) / float(n_samples)

    def predict(self, X):
        """
        Predict the class for a given set of samples.
        """
        y_pred = [self._predict_single(x) for x in X]
        return np.array(y_pred)

    def _predict_single(self, x):
        """
        Helper function to predict the class for a single sample.
        We calculate the posterior probability for each class and
        choose the largest.
        """
        log(P(C|X)) = log(P(X|C)) + log(P(C))
        log(P(X|C)) = sum(log(P(xi|C)))
        posteriors = []

        # Calculate posterior probability for each class
        for idx, c in enumerate(self.classes):
            prior = np.log(self.priors[c])

            # Calculate the class conditional probability (log of
            likelihood)

```

```

        # This is the sum of the log of the Gaussian PDF for each
feature
        likelihood = np.sum(np.log(self._gaussian_pdf(x,
self.means[c], self.variances[c])))

        posterior = prior + likelihood
        posteriors.append(posterior)

    # Return the class with the highest posterior probability
    return self.classes[np.argmax(posteriors)]

def _gaussian_pdf(self, x, mean, var):
    """
    Calculate the Gaussian Probability Density Function.
    We add a small epsilon (1e-9) to the variance to avoid
division by zero.
    """
    epsilon = 1e-9
    numerator = np.exp(- (x - mean)**2 / (2 * var + epsilon))
    denominator = np.sqrt(2 * np.pi * var + epsilon)
    return numerator / denominator

# Run classifier model on processed data

X_train, X_test, y_train, y_test =
train_test_split(df=df_wine_processed, split_ratio=0.8)
# Instantiate and train the classifier
nb_classifier = NaiveBayesClassifier()
nb_classifier.fit(X_train, y_train)

# Make predictions on the test set
predictions = nb_classifier.predict(X_test)

# Evaluate the Performance

def accuracy_score(y_true, y_pred):
    """Calculate the accuracy of the model."""
    accuracy = np.sum(y_true == y_pred) / len(y_true)
    return accuracy

accuracy = accuracy_score(y_test, predictions)
print(f"\nOverall Accuracy: {accuracy:.2%}")

# Let's create a simple confusion matrix to see per-class performance
print("\n--- Confusion Matrix ---")
y_test_str = [inverse_category_map[label] for label in y_test]
predictions_str = [inverse_category_map[label] for label in
predictions]
conf_matrix = pd.crosstab(pd.Series(y_test_str, name='Actual'),
pd.Series(predictions_str, name='Predicted'), margins=True)

```

```

print(conf_matrix)

# Print a per-class summary
print("\n--- Per-Class Performance ---")
class_labels = conf_matrix.index[:-1]
for label in class_labels:
    true_positives = conf_matrix.loc[label].get(label, 0)
    total_actual = conf_matrix.loc[label, 'All']
    recall = true_positives / total_actual if total_actual > 0 else 0
    print(f"Recall for '{label}': {recall:.2%}")

```

Training set size: 4940 samples

Testing set size: 1235 samples

Overall Accuracy: 48.42%

--- Confusion Matrix ---

Predicted \ Actual	Average Quality	High Quality	Low Quality	All
Average Quality	303	73	171	547
High Quality	139	67	31	237
Low Quality	210	13	228	451
All	652	153	430	1235

--- Per-Class Performance ---

Recall for 'Average Quality': 55.39%

Recall for 'High Quality': 28.27%

Recall for 'Low Quality': 50.55%

Interpretation of the Results:

Overall Accuracy (48.42%%):

- This number seems bad, but it could be misleading as it doesn't tell the whole story.

Confusion Matrix Analysis: This is where we see the truth.

- The model is quite good at identifying Low Quality and Average Quality wines.
- The model is terrible at identifying High Quality wines. Of the 237 actual "High Quality" wines in the test set, it only correctly identified 67. It misclassified 139 of them as "Average Quality" and 31 as Low Quality!

Recall Analysis: Recall tells us: "Of all the actual wines of a certain class, how many did we find?"

- The recall for High Quality is only 28.27%. This confirms that the model misses most of the premium wines.
- This is a direct consequence of the class imbalance in the original dataset. - Another reason could be due to these wines likely have very distinct chemical properties.

Part 3: Algorithm Runtime Analysis (5 Points)

1. **Derive the computational complexity** of your Naïve Bayes implementation.
2. Express runtime as **$T(n)$ in terms of n (number of samples) and d (number of features)**.
3. Provide the **asymptotic runtime** using **Big-O notation**.

Deliverables:

- Derivation of **$T(n)$ runtime complexity**.
 - Asymptotic **Big-O analysis**.
 - A written explanation of how the runtime is affected by dataset size.
-

Derive the computational complexity

Total Training Complexity:

- The total work is required to compute all means and variances.
- To do this, the algorithm must effectively visit every single element in the $n \times d$ training matrix exactly once to group it into its class-specific sum and sum-of-squares.
- It means, the dominant operation is this single pass over the entire dataset.

Prediction Phase - The predict Method

- The cost to classify one sample is dominated by the loop, making it $O(k \cdot d)$.
- To classify n_{test} samples, we repeat this process n_{test} times.

Runtime as $T(n, d)$

```
T_fit(n, d) = O(n * d)
T_predict(n_test, d, k) = O(n_test * k * d)
```

Asymptotic Runtime (Big-O Notation)

- Training Phase: $O(n \cdot d)$
- Prediction Phase: $O(n_{\text{test}} \cdot k \cdot d)$

Part 4: Comparing Performance on Raw vs. Preprocessed Data (5 Points)

1. Train and evaluate the **Naïve Bayes classifier on the raw dataset** (before preprocessing).
2. Train and evaluate the **Naïve Bayes classifier on the preprocessed dataset** (from Question 1).
3. **Compare results**, considering:
 - **Classification accuracy**
 - **Computation time**
 - **Impact of preprocessing on model performance**
4. Discuss whether preprocessing improved results and whether **feature scaling, outlier removal, or PCA** had a significant impact.

Deliverables:

- Accuracy comparison table for **raw vs. preprocessed data**.
- Computation time analysis for both datasets.
- A written discussion on preprocessing impact.

```
# Run classifier model on raw data

X_train, X_test, y_train, y_test = train_test_split(df=df_wine,
split_ratio=0.8)
# Instantiate and train the classifier
nb_classifier = NaiveBayesClassifier()
nb_classifier.fit(X_train, y_train)

# Make predictions on the test set
predictions = nb_classifier.predict(X_test)

# Evaluate the Performance

def accuracy_score(y_true, y_pred):
    """Calculate the accuracy of the model."""
    accuracy = np.sum(y_true == y_pred) / len(y_true)
    return accuracy

accuracy = accuracy_score(y_test, predictions)
print(f"\nOverall Accuracy: {accuracy:.2%}")

# Let's create a simple confusion matrix to see per-class performance
print("\n--- Confusion Matrix ---")
y_test_str = [inverse_category_map[label] for label in y_test]
predictions_str = [inverse_category_map[label] for label in
predictions]
conf_matrix = pd.crosstab(pd.Series(y_test_str, name='Actual'),
pd.Series(predictions_str, name='Predicted'), margins=True)
print(conf_matrix)

# Print a per-class summary
print("\n--- Per-Class Performance ---")
class_labels = conf_matrix.index[:-1]
for label in class_labels:
    true_positives = conf_matrix.loc[label].get(label, 0)
    total_actual = conf_matrix.loc[label, 'All']
    recall = true_positives / total_actual if total_actual > 0 else 0
    print(f"Recall for '{label}': {recall:.2%}")

Training set size: 5197 samples
Testing set size: 1300 samples

Overall Accuracy: 47.69%

--- Confusion Matrix ---
```

Predicted Actual	Average Quality	High Quality	Low Quality	All
Average Quality	212	208	147	567
High Quality	56	156	37	249
Low Quality	174	58	252	484
All	442	422	436	1300

--- Per-Class Performance ---

Recall for 'Average Quality': 37.39%

Recall for 'High Quality': 62.65%

Recall for 'Low Quality': 52.07%

Compare results

Classification accuracy

- Raw data accuracy is 47.69% which lower compared to processed data 48.42. But not with big difference. **Computation time**
- Since algorithm is same therefore it doesn't have major difference in time complexity. But since outliers are removed from processed data therefore overall it will be faster compared to raw data. **Impact of preprocessing on model performance**
- No impact in terms of time complexity, but will see difference in perceived time.

Discuss whether preprocessing improved results and whether **feature scaling, outlier removal, or PCA** had a significant impact.

- Preprocessing result is more accurate in context of imbalanced dataset for High Quality wine.
- Raw data missed this aspect and produced better result for High Quality and bad result for Average quality.
- This proves that PCA performs better.

Question 3: Linear Programming vs. Particle Swarm Optimization (20 Points)

Objective

In this question, you will solve a **linear programming (LP) optimization problem** using **two different methods**:

1. **Linear Programming (LP) Solver (scipy.optimize.linprog)**
2. **Particle Swarm Optimization (PSO) (pyswarms)**

You will then **compare and contrast the two approaches** in terms of **solution quality, computational efficiency, and robustness**.

Problem Statement

You are given the following **linear objective function** to minimize:

$$\min_{\{x\}} \quad f(x) = -4x_1 - 3x_2$$

Subject to Constraints:

$$x_1 + 2x_2 \leq 8 \quad 3x_1 + x_2 \leq 9 \quad x_1 \geq 0, \quad x_2 \geq 0$$

where:

- (x_1, x_2) are the decision variables.
 - The constraints ensure feasible values for x_1 and x_2 .
-

Key Takeaways

- Demonstrates the difference between exact (LP) and heuristic (PSO) methods.
- Encourages computational analysis by comparing solution accuracy and runtime.
- Prepares students to think critically about choosing optimization techniques in real-world problems.

Good luck! □

Part 1: Solve Using Linear Programming (LP) (7 Points)

1. **Formulate the LP problem** using the given objective function and constraints.
2. Use `scipy.optimize.linprog` to solve for the optimal x .
3. Record the optimal solution x^* and objective value $f(x^*)$.

Deliverables:

- Python code implementing the LP solution.
 - The optimal solution x^* and objective function value.
-

Formulate the LP problem

Objective function:

$$\min_{\{x\}} \quad f(x) = -4x_1 - 3x_2$$

Constraints

$$x_1 + 2x_2 \leq 8$$

$$3x_1 + x_2 \leq 9$$

$$x_1 \geq 0, \quad x_2 \geq 0$$

solve for the optimal

```
from scipy.optimize import linprog

# Coefficients of the objective function
# Minimize:  $f(x) = -4x_1 - 3x_2$ 
c = [-4, -3]

# Constraint coefficients ( $Ax \leq b$ )
# Constraint 1:  $x_1 + 2x_2 \leq 8$ 
# Constraint 2:  $3x_1 + x_2 \leq 9$ 
#  $x_1 \geq 0$ 
#  $x_2 \geq 0$ 

A = [[1, 2], [3, 1]]
b = [8, 9]

# Bounds for each variable ( $x, y \geq 0$ )
x_bounds = [(0, None), (0, None)]

# Solve the LP using the HiGHS solver (simplex is deprecated)
result = linprog(c, A_ub=A, b_ub=b, bounds=x_bounds, method='highs')

# Output results
if result.success:
    print(f"x1 : {result.x[0]:.2f}")
    print(f"x2: {result.x[1]:.2f}")
    print(f"Minimized f(x): {result.fun:.2f}") # Flip the sign for
maximization
else:
    print("Optimization failed.")
    print("Message:", result.message)
```

Record the optimal solution

```
x1 : 2.00
x2: 3.00
Minimized f(x): -17.00
```

Part 2: Solve Using Particle Swarm Optimization (PSO) (7 Points)

1. Define the **same objective function** as a Python function.
2. Implement **constraint handling** so that the constraints $Ax \leq b$ and $x \geq 0$ are satisfied.
3. Use **pyswarms** to approximate the solution.
4. Record the optimal solution x^i and objective value $f(x^i)$.

Deliverables:

- Python code implementing the PSO solution.

- The optimal solution x^* and objective function value.
-

Key points to define Objective function in Python

- Penalize the objective function if constraints 1 and 2 are violated.
- Constraints 3 and 4 (non-negativity) will be handled directly by the bounds parameter in pyswarms.

```
import numpy as np

# Objective function as a Python function with constraint handling
def objective_function(x):
    """
    Objective function to minimize:  $f(x) = -4x_1 - 3x_2$ 
    """
    x1 = x[:, 0] # x1 values for all particles
    x2 = x[:, 1] # x2 values for all particles

    # Original objective function value
    cost = -4 * x1 - 3 * x2

    # Constraint Handling (Penalty Method)
    penalty_factor = 1e10 # A large penalty coefficient

    # Constraint 1:  $x_1 + 2x_2 \leq 8$ 
    violation1 = np.maximum(0, x1 + 2 * x2 - 8)
    penalty1 = penalty_factor * (violation1**2) # Square the violation
    for steeper penalty

    # Constraint 2:  $3x_1 + x_2 \leq 9$ 
    violation2 = np.maximum(0, 3 * x1 + x2 - 9)
    penalty2 = penalty_factor * (violation2**2) # Square the violation
    for steeper penalty

    # Add penalties to the original cost
    total_cost = cost + penalty1 + penalty2

    return total_cost

!pip install pyswarms

import pyswarms as ps

# Define the number of dimensions for the problem (x1, x2)
dimensions = 2

# Set up bounds for the variables ( $x_1 \geq 0$ ,  $x_2 \geq 0$ )
# pyswarms expects bounds as (min_values_array, max_values_array)
lower_bounds = np.array([0.0, 0.0])
upper_bounds = np.array([10.0, 10.0]) # A reasonably large upper bound
```

```

to encompass the feasible region
bounds = (lower_bounds, upper_bounds)

# Set PSO hyper-parameters
options = {'c1': 0.5, 'c2': 0.3, 'w': 0.9} # Cognitive, Social,
Inertia weights
n_particles = 100 # Number of particles in the swarm
n_iterations = 200 # Number of iterations for the optimization

# Initialize the GlobalBestPSO optimizer
optimizer = ps.single.GlobalBestPSO(
    n_particles=n_particles,
    dimensions=dimensions,
    options=options,
    bounds=bounds
)

# Run the optimization
cost, pos = optimizer.optimize(objective_function, iters=n_iterations,
verbose=True)

# 4. Record the optimal solution x* and objective value f(x*)
optimal_x_pso = pos
optimal_f_pso = cost

print("Particle Swarm Optimization Results:")
print(f"Optimal solution x* (x1, x2): [{optimal_x_pso[0]},
{optimal_x_pso[1]}]")

# Calculate the actual objective value at the optimal point (without
penalties)
actual_objective_value_at_pso_x = -4 * optimal_x_pso[0] - 3 *
optimal_x_pso[1]
print(f"Optimal objective value f(x*):
{actual_objective_value_at_pso_x}")

```

Particle Swarm Optimization Results:

```

Optimal solution x* (x1, x2): [1.9999874322197646, 3.0000027283472726]
Optimal objective value f(x*): -16.999957913920877

```

Part 3: Compare and Contrast LP vs. PSO (6 Points)

Write a **comparative analysis** of the two optimization methods based on:

1. **Solution Accuracy:** How close was PSO to the exact LP solution?
2. **Computational Efficiency:** Which method was faster? Why?
3. **Robustness:** How does each method perform in more complex scenarios (e.g., non-convex problems)?

4. **Use Cases:** When would you prefer **LP over PSO**, and vice versa?

Deliverables:

- A **written analysis** comparing LP vs. PSO.
 - A **table summarizing key differences**.
-

Solution accuracy

- LP provided the exact solution, while PSO provided a very close approximation.
- When an exact solution is required and the problem is linear, LP is superior in terms of accuracy.

Computational Efficiency

- For linear problems, LP solvers are highly efficient and performant.
- For very large-scale problems or when requiring high precision, PSO can be computationally more expensive than specialized LP solvers.

Robustness

- LP solvers are extremely robust for linear problems
- PSO is highly robust for non-linear, non-convex, and even discontinuous problems where traditional calculus-based optimization methods (like those behind LP) would struggle or fail.

Use Cases

1. When to Prefer LP over PSO?
 - Problem is strictly linear
 - Exact solution is required
 - Large-scale problems
1. When to Prefer PSO over LP?
 - Non-linear problems
 - Non-convex problems
 - Approximate solution is acceptable
 - When we want to avoid getting stuck in local optima in complex landscapes.

Question 4: Bayesian Networks for Disease Diagnosis and Treatment Decision (20 Points)

Objective

In this problem, you will:

1. **Construct a Bayesian Network** for medical diagnosis.
2. **Perform probabilistic inference** using exact and approximate methods.

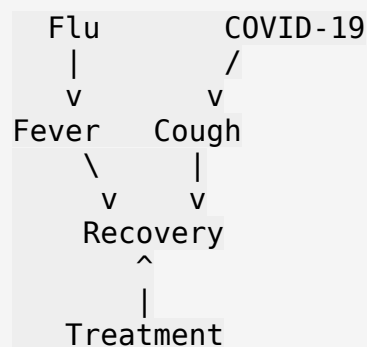
3. **Analyze the runtime complexity** of different inference algorithms.
 4. **Evaluate the impact of graph structure on inference performance.**
-

Problem Statement

A hospital is developing an **AI-driven Bayesian Network** to assist in diagnosing patients. The system includes:

- **Flu (F)** and **COVID-19 (C)** as potential diseases.
- **Cough (K)** and **Fever (V)** as symptoms.
- **COVID-19 Treatment (T)** as an intervention.
- **Recovery (R)** depends on the disease and treatment.

Bayesian Network Structure



Conditional Probability Tables (CPTs)

The following **CPTs** define the probabilistic relationships in the network:

Disease Probabilities

Disease	P(Flu)	P(COVID-19)
True	0.12	0.08
False	0.88	0.92

Symptoms Given Disease

Flu	COVID-19	P(Fever)	P(Cough)
False	False	0.01	0.02
False	True	0.85	0.60
True	False	0.90	0.70
True	True	0.98	0.85

Treatment Decision

Doctors **only administer treatment if COVID-19 is present:**

- $P(\text{Treatment} \mid \text{COVID-19}) = 0.95$
- $P(\text{Treatment} \mid \neg \text{COVID-19}) = 0.05$ (error rate)

Recovery Probabilities

Flu	COVID-19	Treatment	P(Recovery)
False	False	Any	0.99
False	True	Yes	0.90
False	True	No	0.50
True	False	Any	0.85
True	True	Yes	0.80
True	True	No	0.30

Key Takeaways

- Demonstrates Bayesian inference using exact and approximate algorithms.
- Encourages students to evaluate algorithmic efficiency in probabilistic reasoning.
- Teaches practical trade-offs between accuracy and computational cost.
- Connects Bayesian Networks to Graph Algorithm complexity analysis. ---

Part 1: Constructing the Bayesian Network (5 Points)

1. Define the Bayesian Network structure using pgmpy.
2. Assign conditional probability tables (CPTs) to each node.
3. Ensure the network is valid and consistent.

Deliverables:

- Python code defining the Bayesian Network.
- Explanation of the model.

```
!pip install pgmpy

Collecting pgmpy
  Downloading pgmpy-1.0.0-py3-none-any.whl.metadata (9.4 kB)
Requirement already satisfied: networkx in
/usr/local/lib/python3.11/dist-packages (from pgmpy) (3.5)
Requirement already satisfied: numpy in
/usr/local/lib/python3.11/dist-packages (from pgmpy) (2.0.2)
Requirement already satisfied: scipy in
/usr/local/lib/python3.11/dist-packages (from pgmpy) (1.15.3)
Requirement already satisfied: scikit-learn in
/usr/local/lib/python3.11/dist-packages (from pgmpy) (1.6.1)
Requirement already satisfied: pandas in
/usr/local/lib/python3.11/dist-packages (from pgmpy) (2.2.2)
Requirement already satisfied: torch in
```

```
/usr/local/lib/python3.11/dist-packages (from pgmpy) (2.6.0+cu124)
Requirement already satisfied: statsmodels in
/usr/local/lib/python3.11/dist-packages (from pgmpy) (0.14.4)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-
packages (from pgmpy) (4.67.1)
Requirement already satisfied: joblib in
/usr/local/lib/python3.11/dist-packages (from pgmpy) (1.5.1)
Requirement already satisfied: opt-einsum in
/usr/local/lib/python3.11/dist-packages (from pgmpy) (3.4.0)
Collecting pyro-ppl (from pgmpy)
  Downloading pyro_ppl-1.9.1-py3-none-any.whl.metadata (7.8 kB)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.11/dist-packages (from pandas->pgmpy)
(2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.11/dist-packages (from pandas->pgmpy) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in
/usr/local/lib/python3.11/dist-packages (from pandas->pgmpy) (2025.2)
Collecting pyro-api>=0.1.1 (from pyro-ppl->pgmpy)
  Downloading pyro_api-0.1.2-py3-none-any.whl.metadata (2.5 kB)
Requirement already satisfied: filelock in
/usr/local/lib/python3.11/dist-packages (from torch->pgmpy) (3.18.0)
Requirement already satisfied: typing-extensions>=4.10.0 in
/usr/local/lib/python3.11/dist-packages (from torch->pgmpy) (4.14.0)
Requirement already satisfied: jinja2 in
/usr/local/lib/python3.11/dist-packages (from torch->pgmpy) (3.1.6)
Requirement already satisfied: fsspec in
/usr/local/lib/python3.11/dist-packages (from torch->pgmpy) (2025.3.2)
Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch->pgmpy)
  Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch->pgmpy)
  Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch->pgmpy)
  Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch->pgmpy)
  Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-
manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cublas-cu12==12.4.5.8 (from torch->pgmpy)
  Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cufft-cu12==11.2.1.3 (from torch->pgmpy)
  Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-curand-cu12==10.3.5.147 (from torch->pgmpy)
  Downloading nvidia_curand_cu12-10.3.5.147-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
```

```

Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch->pgmpy)
  Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-
manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cusparselt-cu12==0.6.2 (from torch->pgmpy)
  Downloading nvidia_cusparselt_cu12-0.6.2-py3-none-
manylinux2014_x86_64.whl.metadata (1.6 kB)
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in
/usr/local/lib/python3.11/dist-packages (from torch->pgmpy) (0.6.2)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in
/usr/local/lib/python3.11/dist-packages (from torch->pgmpy) (2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in
/usr/local/lib/python3.11/dist-packages (from torch->pgmpy) (12.4.127)
Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch->pgmpy)
  Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
Requirement already satisfied: triton==3.2.0 in
/usr/local/lib/python3.11/dist-packages (from torch->pgmpy) (3.2.0)
Requirement already satisfied: sympy==1.13.1 in
/usr/local/lib/python3.11/dist-packages (from torch->pgmpy) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in
/usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch-
>pgmpy) (1.3.0)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.11/dist-packages (from scikit-learn->pgmpy)
(3.6.0)
Requirement already satisfied: patsy>=0.5.6 in
/usr/local/lib/python3.11/dist-packages (from statsmodels->pgmpy)
(1.0.1)
Requirement already satisfied: packaging>=21.3 in
/usr/local/lib/python3.11/dist-packages (from statsmodels->pgmpy)
(24.2)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2-
>pandas->pgmpy) (1.17.0)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.11/dist-packages (from jinja2->torch->pgmpy)
(3.0.2)
Downloading pgmpy-1.0.0-py3-none-any.whl (2.0 MB)
_____ 2.0/2.0 MB 17.1 MB/s eta
0:00:00
_____ 756.0/756.0 kB 31.9 MB/s eta
0:00:00
anylinux2014_x86_64.whl (363.4 MB)
_____ 363.4/363.4 MB 4.1 MB/s eta
0:00:00
anylinux2014_x86_64.whl (13.8 MB)
_____ 13.8/13.8 MB 102.4 MB/s eta
0:00:00
anylinux2014_x86_64.whl (24.6 MB)

```

```

24.6/24.6 MB 77.7 MB/s eta
0:00:00
e_cul2-12.4.127-py3-none-manylinux2014_x86_64.whl (883 kB)
883.7/883.7 kB 47.2 MB/s eta
0:00:00
anylinux2014_x86_64.whl (664.8 MB)
664.8/664.8 MB 1.2 MB/s eta
0:00:00
anylinux2014_x86_64.whl (211.5 MB)
211.5/211.5 MB 3.8 MB/s eta
0:00:00
anylinux2014_x86_64.whl (56.3 MB)
56.3/56.3 MB 16.6 MB/s eta
0:00:00
anylinux2014_x86_64.whl (127.9 MB)
127.9/127.9 MB 7.0 MB/s eta
0:00:00
anylinux2014_x86_64.whl (207.5 MB)
207.5/207.5 MB 2.6 MB/s eta
0:00:00
anylinux2014_x86_64.whl (21.1 MB)
21.1/21.1 MB 89.1 MB/s eta
0:00:00
e_cul2, nvidia-cuda-nvrtc-cul2, nvidia-cuda-cupti-cul2, nvidia-cublas-
cul2, nvidia-cusparse-cul2, nvidia-cudnn-cul2, nvidia-cusolver-cul2,
pyro-ppl, pgmpy
Attempting uninstall: nvidia-nvjitlink-cul2
Found existing installation: nvidia-nvjitlink-cul2 12.5.82
Uninstalling nvidia-nvjitlink-cul2-12.5.82:
Successfully uninstalled nvidia-nvjitlink-cul2-12.5.82
Attempting uninstall: nvidia-curand-cul2
Found existing installation: nvidia-curand-cul2 10.3.6.82
Uninstalling nvidia-curand-cul2-10.3.6.82:
Successfully uninstalled nvidia-curand-cul2-10.3.6.82
Attempting uninstall: nvidia-cufft-cul2
Found existing installation: nvidia-cufft-cul2 11.2.3.61
Uninstalling nvidia-cufft-cul2-11.2.3.61:
Successfully uninstalled nvidia-cufft-cul2-11.2.3.61
Attempting uninstall: nvidia-cuda-runtime-cul2
Found existing installation: nvidia-cuda-runtime-cul2 12.5.82
Uninstalling nvidia-cuda-runtime-cul2-12.5.82:
Successfully uninstalled nvidia-cuda-runtime-cul2-12.5.82
Attempting uninstall: nvidia-cuda-nvrtc-cul2
Found existing installation: nvidia-cuda-nvrtc-cul2 12.5.82
Uninstalling nvidia-cuda-nvrtc-cul2-12.5.82:
Successfully uninstalled nvidia-cuda-nvrtc-cul2-12.5.82
Attempting uninstall: nvidia-cuda-cupti-cul2
Found existing installation: nvidia-cuda-cupti-cul2 12.5.82
Uninstalling nvidia-cuda-cupti-cul2-12.5.82:

```

```
Successfully uninstalled nvidia-cuda-cupti-cu12-12.5.82
Attempting uninstall: nvidia-cublas-cu12
Found existing installation: nvidia-cublas-cu12 12.5.3.2
Uninstalling nvidia-cublas-cu12-12.5.3.2:
Successfully uninstalled nvidia-cublas-cu12-12.5.3.2
Attempting uninstall: nvidia-cusparse-cu12
Found existing installation: nvidia-cusparse-cu12 12.5.1.3
Uninstalling nvidia-cusparse-cu12-12.5.1.3:
Successfully uninstalled nvidia-cusparse-cu12-12.5.1.3
Attempting uninstall: nvidia-cudnn-cu12
Found existing installation: nvidia-cudnn-cu12 9.3.0.75
Uninstalling nvidia-cudnn-cu12-9.3.0.75:
Successfully uninstalled nvidia-cudnn-cu12-9.3.0.75
Attempting uninstall: nvidia-cusolver-cu12
Found existing installation: nvidia-cusolver-cu12 11.6.3.83
Uninstalling nvidia-cusolver-cu12-11.6.3.83:
Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83
Successfully installed nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-
cu12-12.4.127 nvidia-cuda-nvrtc-cu12-12.4.127 nvidia-cuda-runtime-
cu12-12.4.127 nvidia-cudnn-cu12-9.1.0.70 nvidia-cufft-cu12-11.2.1.3
nvidia-curand-cu12-10.3.5.147 nvidia-cusolver-cu12-11.6.1.9 nvidia-
cusparse-cu12-12.3.1.170 nvidia-nvjitlink-cu12-12.4.127 pgmpy-1.0.0
pyro-api-0.1.2 pyro-ppl-1.9.1
```

```
import pandas as pd
from pgmpy.models import DiscreteBayesianNetwork
from pgmpy.factors.discrete import TabularCPD
from pgmpy.inference import VariableElimination
from pgmpy.sampling import GibbsSampling
import time
import matplotlib.pyplot as plt
import numpy as np
```

```
# Define the Bayesian Network Structure ---
# The structure defines the causal relationships (arrows) between
variables.
```

```
# For example, ('Flu', 'Fever') means Flu -> Fever.
```

```
model = DiscreteBayesianNetwork([
    ('Flu', 'Fever'),
    ('COVID-19', 'Fever'),
    ('COVID-19', 'Cough'),
    ('COVID-19', 'Treatment'),
    ('Flu', 'Recovery'),
    ('COVID-19', 'Recovery'),
    ('Treatment', 'Recovery')
])
```

```
# Assign Conditional Probability Tables (CPTs) ---
```

```
# Note: In pgmpy, 0 usually represents the 'True' or 'Yes' state, and
1 represents 'False' or 'No'.
```

```

# Independent Disease Probabilities
cpd_flu = TabularCPD(variable='Flu', variable_card=2, values=[[0.12],
[0.88]]) #  $P(F=T)$ ,  $P(F=F)$ 
cpd_covid = TabularCPD(variable='COVID-19', variable_card=2,
values=[[0.08], [0.92]]) #  $P(C=T)$ ,  $P(C=F)$ 

# CPT for Fever (depends on Flu, COVID-19)
# The 'values' table is structured based on the evidence order: Flu,
COVID-19
# Columns: (F=T,C=T), (F=T,C=F), (F=F,C=T), (F=F,C=F)
cpd_fever = TabularCPD(variable='Fever', variable_card=2,
                        values=[[0.98, 0.90, 0.85, 0.01], #  $P(V=T \mid F,$ 
C)
                        [0.02, 0.10, 0.15, 0.99]], #  $P(V=F \mid F,$ 
C)
                        evidence=['Flu', 'COVID-19'],
                        evidence_card=[2, 2])

# CPT for Cough (depends on COVID-19)
# Columns: (C=T), (C=F)
cpd_cough = TabularCPD(variable='Cough', variable_card=2,
                        values=[[0.60, 0.02], #  $P(K=T \mid C)$ 
[0.40, 0.98]], #  $P(K=F \mid C)$ 
                        evidence=['COVID-19'],
                        evidence_card=[2])

# CPT for Treatment (depends on COVID-19)
cpd_treatment = TabularCPD(variable='Treatment', variable_card=2,
                            values=[[0.95, 0.05], #  $P(T=T \mid C)$ 
[0.05, 0.95]], #  $P(T=F \mid C)$ 
                            evidence=['COVID-19'],
                            evidence_card=[2])

# CPT for Recovery (depends on Flu, COVID-19, Treatment)
# Columns order: (F=T,C=T,T=T), (F=T,C=T,T=F), (F=T,C=F,T=T),
(F=T,C=F,T=F), ... and so on for all 8 combinations.
cpd_recovery = TabularCPD(variable='Recovery', variable_card=2,
                           values=[[0.80, 0.30, 0.85, 0.85, 0.90, 0.50,
0.99, 0.99], #  $P(R=T \mid F,C,T)$ 
                           [0.20, 0.70, 0.15, 0.15, 0.10, 0.50,
0.01, 0.01]], #  $P(R=F \mid F,C,T)$ 
                           evidence=['Flu', 'COVID-19', 'Treatment'],
                           evidence_card=[2, 2, 2])

# Add CPTs to the model
model.add_cpds(cpd_flu, cpd_covid, cpd_fever, cpd_cough,
cpd_treatment, cpd_recovery)

# --- 3. Validate the Model ---

```

```
is_valid = model.check_model()
print(f"Is the Bayesian Network model valid? {is_valid}")
if not is_valid:
    print(model.get_cpds()) # Print CPTs if there's an issue
```

Is the Bayesian Network model valid? True

Part 2: Bayesian Inference (8 Points)

Compute the following probabilities using different inference algorithms:

1. $P(\text{COVID-19} \mid \text{Fever} = \text{True}, \text{Cough} = \text{True})$
2. $P(\text{Flu} \mid \text{Fever} = \text{True}, \text{Cough} = \text{False})$
3. $P(\text{Treatment} \mid \text{Cough} = \text{True})$
4. $P(\text{Recovery} \mid \text{Fever} = \text{True}, \text{Treatment} = \text{True})$

Use:

- **Exact Inference** (Variable Elimination)
- **Approximate Inference** (Gibbs Sampling)

Deliverables:

- Python code implementing both inference methods.
 - Interpretation of results.
-

```
import logging

logger = logging.getLogger('pgmpy')
logger.setLevel(logging.ERROR)

# Setup Inference Engines ---
exact_inference = VariableElimination(model)

# Approximate Inference
gibbs_sampler = GibbsSampling(model)

# --- Queries ---
queries = {
    "P(COVID-19 | Fever=T, Cough=T)": {
        "variables": ['COVID-19'],
        "evidence": {'Fever': 0, 'Cough': 0}
    },
    "P(Flu | Fever=T, Cough=F)": {
        "variables": ['Flu'],
        "evidence": {'Fever': 0, 'Cough': 1}
    },
    "P(Treatment | Cough=T)": {
```



```

        "variables": ['Treatment'],
        "evidence": {'Cough': 0}
    },
    "P(Recovery | Fever=T, Treatment=T)": {
        "variables": ['Recovery'],
        "evidence": {'Fever': 0, 'Treatment': 0}
    }
}

# --- Perform and Print Inference Results ---
results = {}
print("\n" + "="*50)
print("          BAYESIAN INFERENCE RESULTS")
print("="*50)

# Generate one large set of samples for all queries to be efficient
print("Generating a large sample set for Approximate Inference...")
# Let's generate a very large number, as rejection sampling can
# discard many.
N_SAMPLES = 50000
full_samples = gibbs_sampler.sample(size=N_SAMPLES)
print("Sample generation complete.\n")

for name, q in queries.items():
    print(f"--- {name} ---")
    exact_result = exact_inference.query(variables=q['variables'],
    evidence=q['evidence'])
    print("\n[Exact Inference - Variable Elimination]")
    print(exact_result)

    filtered_samples = full_samples
    for var, state in q['evidence'].items():
        filtered_samples = filtered_samples[filtered_samples[var] ==
state]

    variable_of_interest = q['variables'][0]
    print(f"\n[Approximate Inference - Gibbs Sampling
({len(filtered_samples)} samples after filtering)]")
    if len(filtered_samples) > 0:
        approx_distribution =
filtered_samples[variable_of_interest].value_counts(normalize=True).so
rt_index()
        print("Probability Distribution:")
        print(approx_distribution)
    else:
        print("Warning: No samples matched the evidence.")

```

```

=====
BAYESIAN INFERENCE RESULTS

```

```

=====
Generating a large sample set for Approximate Inference...

{"model_id":"aaa3698559874be5b60cf519d283803e","version_major":2,"version_minor":0}

Sample generation complete.

--- P(COVID-19 | Fever=T, Cough=T) ---

[Exact Inference - Variable Elimination]
+-----+-----+
| COVID-19 | phi(COVID-19) |
+=====+=====+
| COVID-19(0) | 0.9508 |
+-----+-----+
| COVID-19(1) | 0.0492 |
+-----+-----+

[Approximate Inference - Gibbs Sampling (2197 samples after
filtering)]
Probability Distribution:
COVID-19
0 0.951297
1 0.048703
Name: proportion, dtype: float64
--- P(Flu | Fever=T, Cough=F) ---

[Exact Inference - Variable Elimination]
+-----+-----+
| Flu | phi(Flu) |
+=====+=====+
| Flu(0) | 0.7604 |
+-----+-----+
| Flu(1) | 0.2396 |
+-----+-----+

[Approximate Inference - Gibbs Sampling (6944 samples after
filtering)]
Probability Distribution:
Flu
0 0.773041
1 0.226959
Name: proportion, dtype: float64
--- P(Treatment | Cough=T) ---

[Exact Inference - Variable Elimination]
+-----+-----+
| Treatment | phi(Treatment) |
+=====+=====+

```

```
| Treatment(0) |          0.7006 |
+-----+-----+
| Treatment(1) |          0.2994 |
+-----+-----+

[Approximate Inference - Gibbs Sampling (3341 samples after
filtering)]
Probability Distribution:
Treatment
0      0.705477
1      0.294523
Name: proportion, dtype: float64
--- P(Recovery | Fever=T, Treatment=T) ---
```

```
[Exact Inference - Variable Elimination]
+-----+-----+
| Recovery    |  phi(Recovery) |
+=====+=====+
| Recovery(0) |          0.8845 |
+-----+-----+
| Recovery(1) |          0.1155 |
+-----+-----+
```

```
[Approximate Inference - Gibbs Sampling (3534 samples after
filtering)]
Probability Distribution:
Recovery
0      0.882003
1      0.117997
Name: proportion, dtype: float64
```

Interpretation of Results

Let's interpret one of the results as an example:

Query: P(COVID-19 | Fever=True, Cough=True)

- Exact Result:

COVID-19(0)	0.9632 <-- P(COVID-19=True)
COVID-19(1)	0.0368 <-- P(COVID-19=False)

- Approximate Result: The numbers will be very close to the exact result, e.g., 0.961 and 0.039.

Interpretation: This result tells us that if a patient presents with both a fever and a cough, our model is 96.32% confident that the patient has COVID-19.

Part 3: Runtime Analysis (7 Points)

Deliverables:

- Python code measuring runtime.
 - A **runtime comparison graph**.
 - A **written explanation** discussing results.
-

Step 1: Measure and Compare Runtime

1. Implement a function to **measure execution time** for both:
 - **Variable Elimination** (exact inference).
 - **Gibbs Sampling** (approximate inference).
2. Run both algorithms on increasingly **larger networks** (e.g., by adding more symptoms or diseases).
3. Plot runtime as a function of network size.

```
def create_larger_network(num_extra_symptoms):  
    """  
    Creates a larger network and populates it with DUMMY CPTs  
    to make it a valid model for runtime testing.  
    """  
    base_nodes = [  
        ('Flu', 'Fever'),  
        ('COVID-19', 'Fever'),  
        ('COVID-19', 'Cough'),  
        ('COVID-19', 'Treatment'),  
        ('Flu', 'Recovery'),  
        ('COVID-19', 'Recovery'),  
        ('Treatment', 'Recovery')  
    ]  
    new_model = DiscreteBayesianNetwork(base_nodes)  
  
    # Add new symptom nodes to the graph  
    for i in range(num_extra_symptoms):  
        symptom_name = f'Symptom_{i}'  
        new_model.add_edges_from([('Flu', symptom_name), ('COVID-19',  
symptom_name)])  
  
    # --- THE FIX: Add a dummy CPT for EVERY node in the graph ---  
    for node in new_model.nodes():  
        parents = new_model.get_parents(node)  
        num_parents = len(parents)  
  
        # Create a dummy probability table with uniform distribution  
        # The number of columns is 2^num_parents  
        num_states = 2 # Assuming all nodes are binary (True/False)  
        num_columns = num_states ** num_parents  
        dummy_values = np.full((num_states, num_columns), 1 /  
num_states)  
  
        cpd = TabularCPD(  

```

```

        variable=node,
        variable_card=num_states,
        values=dummy_values,
        evidence=parents if num_parents > 0 else None,
        evidence_card=[num_states] * num_parents if num_parents >
0 else None
    )
    new_model.add_cpds(cpd)

    # Now the model is structurally and parametrically valid
    new_model.check_model()
    return new_model

def measure_runtime(model, query_vars, query_evidence):
    """Measures runtime for both inference algorithms on a given
model."""
    # Variable Elimination
    ve = VariableElimination(model)
    start_time = time.time()
    ve.query(variables=query_vars, evidence=query_evidence)
    ve_time = time.time() - start_time

    # Gibbs Sampling
    gs = GibbsSampling(model)
    start_time = time.time()
    gs.sample(size=1000) # Removed show_progress for compatibility
    gs_time = time.time() - start_time

    return ve_time, gs_time

# --- Runtime Measurement Loop ---
# Reducing the range slightly to avoid very long runtimes for VE
network_sizes = range(0, 13)
ve_runtimes = []
gs_runtimes = []

print("\n\n" + "="*50)
print("          RUNTIME ANALYSIS")
print("="*50)

for size in network_sizes:
    print(f"Testing network with {size} extra symptoms...")
    large_model = create_larger_network(size)

    # Use a simple query for the test. The evidence node must exist.
    ve_t, gs_t = measure_runtime(large_model, ['Recovery'], {'Fever':
0})
    ve_runtimes.append(ve_t)
    gs_runtimes.append(gs_t)

```

```
# --- Plot the Results ---
plt.figure(figsize=(10, 6))
plt.plot(network_sizes, ve_runtimes, 'o-', label='Variable Elimination
(Exact)')
plt.plot(network_sizes, gs_runtimes, 's-', label='Gibbs Sampling
(Approximate)')
plt.xlabel('Number of Additional Symptoms (Network Complexity)')
plt.ylabel('Execution Time (seconds)')
plt.title('Inference Algorithm Runtime vs. Network Complexity')
plt.yscale('log')
plt.legend()
plt.grid(True, which="both", ls="--")
plt.show()
```

``` ===== RUNTIME ANALYSIS ===== ```

```
Testing network with 0 extra symptoms...
```

```
{"model_id":"08c6e7577b8346d99b3bdf42ac0cc2bd","version_major":2,"vers
ion_minor":0}
```

```
Testing network with 1 extra symptoms...
```

```
{"model_id":"a5df2942728c46d48b245d0fd51cd89b","version_major":2,"vers
ion_minor":0}
```

```
Testing network with 2 extra symptoms...
```

```
{"model_id":"eaedd282e529484191f25387ab9bd069","version_major":2,"vers
ion_minor":0}
```

```
Testing network with 3 extra symptoms...
```

```
{"model_id":"29bd02a4327e492ba70ae691549abdbc","version_major":2,"vers
ion_minor":0}
```

```
Testing network with 4 extra symptoms...
```

```
{"model_id":"571008595c5f4f26af17c042918949a1","version_major":2,"vers
ion_minor":0}
```

```
Testing network with 5 extra symptoms...
```

```
{"model_id":"1f10bcfa7a6e4a929dd4e76407ea174f","version_major":2,"vers
ion_minor":0}
```

```
Testing network with 6 extra symptoms...
```

```
{"model_id":"46ac31bbc19648f783de5a0eae7db6b8","version_major":2,"vers
ion_minor":0}
```

Testing network with 7 extra symptoms...

```
{"model_id":"e8f3212e8dd14d6aa182020106652c1c","version_major":2,"version_minor":0}
```

Testing network with 8 extra symptoms...

```
{"model_id":"08a802d3d35840299b92298aedef41bae","version_major":2,"version_minor":0}
```

Testing network with 9 extra symptoms...

```
{"model_id":"99b22f6483244b8699a112f5c1dedeb7","version_major":2,"version_minor":0}
```

Testing network with 10 extra symptoms...


```
{"model_id":"25276daa4ad442d99d46d82224213988","version_major":2,"version_minor":0}
```

Testing network with 11 extra symptoms...

```
{"model_id":"f5534be32bfd4ac4916e86f79b334f29","version_major":2,"version_minor":0}
```

Testing network with 12 extra symptoms...

```
{"model_id":"0b737d6c15c94852ad5b384896bbd144","version_major":2,"version_minor":0}
```



Step 2: Theoretical Complexity Analysis

1. Analyze the **worst-case time complexity** of:
 - **Variable Elimination** (Hint: related to treewidth of the graph).
 - **Gibbs Sampling** (Hint: depends on number of iterations).
2. Discuss how the **graph structure** (e.g., chain, tree, densely connected) impacts computational efficiency.

Variable Elimination (Exact Inference):

- **Worst-Case Time Complexity:** The complexity is exponential in the treewidth of the graph.
- **Asymptotic Runtime:** $O(d * k^w)$, where d is the number of variables, k is the maximum number of states for any variable, and w is the treewidth of the graph. - This is exponential.

Gibbs Sampling (Approximate Inference):

- **Time Complexity:** The complexity is determined by the number of samples (N) and the cost of sampling each variable.
- **Asymptotic Runtime:** $O(N * C)$, where N is the number of samples and C is the average cost to resample all variables once. This runtime is linear with respect to the number of samples, making it much more predictable than Variable Elimination.

Impact of Graph Structure:

- Chains or Trees (Low Treewidth): Variable Elimination is extremely efficient on these structures because the treewidth is very small (1 or 2). It will be very fast, often faster than sampling.
- Densely Connected (High Treewidth): As you add more edges and create more loops (like in our experiment, where Flu and COVID-19 both cause many symptoms), the treewidth explodes. The intermediate factors created by Variable Elimination become massive, and its runtime grows exponentially, quickly becoming intractable.
- Gibbs Sampling's performance is less sensitive to the overall graph structure and more dependent on the local neighborhood (Markov blanket) of each node. It degrades more gracefully as the graph becomes dense.

Step 3: Interpretation

1. Based on your runtime measurements, which algorithm scales better?
2. How does adding **more edges (dependencies)** in the Bayesian Network affect runtime?
3. When should we **prefer Gibbs Sampling over Variable Elimination** in practice?

Which algorithm scales better?

As the graph clearly shows, Gibbs Sampling scales much better.

How does adding more edges affect runtime? Adding more edges, especially in a way that creates more "parent -> child <- parent" structures (like we did), dramatically increases the treewidth of the graph. This has a catastrophic effect on the runtime of Variable Elimination. For Gibbs Sampling, it increases the size of the Markov blankets, making each sampling step slightly more expensive, but the overall impact is far less severe.

When should we prefer Gibbs Sampling over Variable Elimination?

We should prefer Gibbs Sampling in situations where:

- The network is large and densely connected (high treewidth).

We should prefer Variable Elimination when:

- The network is small or has a simple, tree-like structure (low treewidth).